

Let's assume that we have a series system with  $m = 3$  components, each of which has an exponentially distributed time-to-failure,

$$T_{ij} \sim \text{exponential}(\lambda_j)$$

for  $j = 1, \dots, m$  and  $i = 1, \dots, n$  where  $\theta = (\lambda_1, \lambda_2, \lambda_3)$  is unknown.

We replicate the masked data in Table 2 from the Guo paper:

```
md <- read.csv("./data.csv")
md$X1 <- ifelse(md$X1 == 0, F, T)
md$X2 <- ifelse(md$X2 == 0, F, T)
md$X3 <- ifelse(md$X3 == 0, F, T)
md
```

##	System.ID	Failure.Time	X1	X2	X3
## 1	1	21	FALSE	TRUE	FALSE
## 2	2	38	TRUE	TRUE	FALSE
## 3	3	54	FALSE	FALSE	TRUE
## 4	4	66	FALSE	FALSE	TRUE
## 5	5	76	TRUE	TRUE	FALSE
## 6	6	78	FALSE	TRUE	TRUE
## 7	7	123	FALSE	FALSE	TRUE
## 8	8	130	TRUE	FALSE	TRUE
## 9	9	152	TRUE	TRUE	TRUE
## 10	10	159	TRUE	FALSE	FALSE
## 11	11	199	FALSE	FALSE	TRUE
## 12	12	201	TRUE	FALSE	FALSE
## 13	13	204	TRUE	FALSE	FALSE
## 14	14	215	FALSE	TRUE	TRUE
## 15	15	218	TRUE	TRUE	FALSE
## 16	16	281	TRUE	FALSE	FALSE
## 17	17	295	FALSE	TRUE	FALSE
## 18	18	310	FALSE	FALSE	TRUE
## 19	19	338	FALSE	FALSE	TRUE
## 20	20	341	FALSE	TRUE	FALSE
## 21	21	354	TRUE	FALSE	FALSE
## 22	22	358	FALSE	TRUE	FALSE
## 23	23	431	TRUE	TRUE	TRUE
## 24	24	457	FALSE	FALSE	TRUE
## 25	25	545	TRUE	TRUE	TRUE
## 26	26	569	FALSE	TRUE	FALSE
## 27	27	677	FALSE	FALSE	TRUE
## 28	28	818	FALSE	TRUE	FALSE
## 29	29	946	FALSE	TRUE	TRUE
## 30	30	1486	TRUE	FALSE	FALSE

We are using the likelihood function given by

$$\begin{aligned}
L(\lambda_1, \lambda_2, \lambda_3) &= \prod_{i=1}^n \sum_{j \in C_i} f_j(t_i; \theta_j) \prod_{\substack{p=1 \\ p \neq j}} R_p(t_i; \theta_p) \\
&= \prod_{i=1}^n \left[ \left\{ \prod_{j=1}^m R_j(t_i; \theta_j) \right\} \left\{ \sum_{k \in C_i} h_k(t_i; \theta_k) \right\} \right]
\end{aligned}$$

and the log-likelihood function given by

$$l(\lambda_1, \lambda_2, \lambda_3) = \sum_{i=1}^n \sum_{j=1}^m \log R_j(t_i; \theta_j) + \sum_{i=1}^n \log \left\{ \sum_{k \in C_i} h_k(t_i; \theta_k) \right\}.$$

The component times-to-failure are exponentially distributed, and thus the hazard and survival functions are respectively given by  $h_j(t; \lambda_j) = \lambda_j$  and  $R_j(t; \lambda_j) = \exp(-\lambda_j t)$ . Making this substitution into the likelihood function obtains the result

$$L(\lambda_1, \lambda_2, \lambda_3) = \prod_{i=1}^n \left( \sum_{j \in C_i} \lambda_j \right) e^{-\left( \sum_{j=1}^m \lambda_j \right) t_i}. \quad (1)$$

and into the log-likelihood function obtains the result

$$\ell(\lambda_1, \lambda_2, \lambda_3) = \sum_{i=1}^n \log \left( \sum_{j \in C_i} \lambda_j \right) - \left( \sum_{j=1}^m \lambda_j \right) \left( \sum_{i=1}^n t_i \right). \quad (2)$$

We model the log-likelihood function in R with:

```
# generator for log-likelihood function
loglike <- function(ttf,cand)
{
  n <- length(ttf)
  function(theta)
  {
    res <- 0
    for (i in 1:n)
    {
      #cat("i=", i, "ttf=", ttf[i], "C=", (1:m)[cand[i,]], "sum=", sum(theta[cand[i,]]), "\n")
      res <- res + log(sum(theta[cand[i,]]))
    }
    return(res - sum(theta) * sum(ttf))
  }
}

mle.grad <- function(l,theta0,eps=1e-3,r=.5)
{
  repeat
  {
    # we use backtracking for an approximate line search
    alpha <- 1
    theta1 <- NULL
    g <- numDeriv::grad(l,theta0)
    repeat {
      theta1 <- theta0 + alpha * g
      if (all(theta1 > 0) && l(theta1) > l(theta0))
        break
      alpha <- r*alpha
    }

    # infinity norm
    if (abs(sum(theta1-theta0)) < eps)
      return(theta1)
    theta0 <- theta1
  }
}
```

So, let's set up the log-likelihood function:

```
ttf <- md$Failure.Time
cand <- as.matrix(md[3:5])
l <- loglike(ttf,cand)
```

Now, we try to solve the MLE using gradient ascent to solve for the zeros of the gradient of the log-likelihood function. We repeat the method many times, using random initial points, and choose the best one found:

```
trials <- 1
theta.hat <- NULL
l.theta.hat <- -Inf
repeat
{
  theta0 <- runif(3,1.5,6.5)
  theta.b <- NULL
  tryCatch(
    {
      theta.b <- mle.grad(l,theta0,1e-6)
      l.theta.b <- l(theta.b)
      if (l.theta.b > l.theta.hat)
      {
        l.theta.hat <- l.theta.b
        theta.hat <- theta.b
      }
      trials <- trials + 1L
    },
    warning = function(w) {},
    error = function(e) {})

  if (trials == 5000L)
    break
}

print(theta.hat)
## [1] 0.0008579827 0.0009880395 0.0011126141
```

So, we see that  $\hat{\theta} = (0.0008579827, 0.0009880395, 0.0011126141)$ .