

Bootstrapping confidence intervals (BCa) of the maximum likelihood estimator of components in a series systems from masked failure data

Alex Towell

Abstract

We estimate the parameters of a series system with Weibull component lifetimes from relatively small samples consisting of right-censored system lifetimes and masked component cause of failure. Under a set of conditions that permit us to ignore how the component cause of failures are masked, we assess the bias and variance of the estimator. Then, we assess the accuracy of the bootstrapped variance and calibration of the confidence intervals of the MLE under a variety of scenarios.

Contents

1	Introduction	2
2	Series System Model	2
2.1	Component Cause of Failure	5
2.2	System and Component Reliabilities	7
3	Likelihood Model for Masked Data	7
3.1	Masked Component Cause of Failure	9
3.2	Right-Censored Data	11
3.3	Identifiability and Convergence Issues	12
4	Maximum Likelihood Estimation	12
5	Bias-Corrected and Accelerated Bootstrap Confidence Intervals	13
5.1	Issues with Resampling from the Observed Data	14
6	Series System with Weibull Components	14
6.1	Reliability	16
6.2	Likelihood Model	16
6.3	Weibull Series System: Homogeneous Shape Parameters	17
7	Simulation Study: Full Model	18
7.1	Data Generating Process	19
7.2	Simulation Scenarios	20
7.3	Scenario: Assessing the Impact of Right-Censoring	20
7.4	Scenario: Assessing the Impact of Sample Size	23
7.5	Scenario: Assessing the Impact of Masking Probability for Component Cause of Failure	25
7.6	Scenario: Assessing the Impact of Changing the Scale Parameter of a Component	25
7.7	Scenario: Assessing the Impact of Changing the Shape Parameter of a Component	27
7.8	Implications and Recommendations	29
8	Simulation Study: Full Model vs Reduced Model	31
8.1	Implications and Recommendations	33
9	Future Work	33

10 Conclusion	34
Appendix A: R Code For Log-likelihood Function	34
Appendix B: R Code For Score Function	35
Appendix C: R Code For Simulation of Scenarios	36
Appendix D: Bernoulli Candidate Set Model	39
Appendix E: Series System Quantile Function	40
Appendix F: Likelihood Ratio Test	40

1 Introduction

Accurately estimating the reliability of individual components in multi-component systems is an important problem in many engineering domains. However, component lifetimes and failure causes are often not directly observable. In a series system, only the system-level failure time may be recorded along with limited information about which component failed. Such *masked* data poses challenges for estimating component reliability.

In this paper, we develop a maximum likelihood approach to estimate component reliability in series systems using right-censored lifetime data and candidate sets that contain the failed component. The key contributions are:

1. Deriving a likelihood model that accounts for right-censoring and masked failure causes through candidate sets. This allows the available masked data to be used for estimation.
2. Validating the accuracy, precision, and robustness of the maximum likelihood estimator through an extensive simulation study under different sample sizes, masking probabilities, and censoring levels.
3. Demonstrating that bootstrapping provides well-calibrated confidence intervals for the MLEs even with small samples.

Together, these contributions provide a statistically rigorous methodology for learning about latent component properties from series system data. The methods are shown to work well even when failure information is significantly masked. This capability expands the range of applications where component reliability can be quantified from limited observations.

The remainder of this paper is organized as follows. First, we detail the series system and masked data models. Next, we present the likelihood construction and maximum likelihood theory. We then describe the bootstrap approach for variance and confidence interval estimation. Finally, we validate the methods through simulation studies under various data scenarios and sample sizes.

2 Series System Model

Consider a system composed of m components arranged in a series configuration. Each component and system has two possible states, functioning or failed. We have n systems whose lifetimes are independent and identically distributed (i.i.d.). The lifetime of the i^{th} system denoted by the random variable T_i . The lifetime of the j^{th} component in the i^{th} system is denoted by the random variable T_{ij} . We assume the component lifetimes in a single system are statistically independent and non-identically distributed. Here, lifetime is defined as the elapsed time from when the new, functioning component (or system) is put into operation until it fails for the first time. A series system fails when any component fails, thus the lifetime of the i^{th} system is given by the component with the shortest lifetime,

$$T_i = \min\{T_{i1}, T_{i2}, \dots, T_{im}\}.$$

There are three particularly important distribution functions in reliability analysis: the reliability function, the probability density function, and the hazard function. The reliability function, $R_{T_i}(t)$, is the probability that the i^{th} system has a lifespan larger than a duration t ,

$$R_{T_i}(t) = \Pr\{T_i > t\} \quad (1)$$

The probability density function (pdf) of T_i is denoted by $f_{T_i}(t)$ and may be defined as

$$f_{T_i}(t) = -\frac{d}{dt}R_{T_i}(t).$$

Next, we introduce the hazard function. The probability that a failure occurs between t and Δt given that no failure occurs before time t is given by

$$\Pr\{T_i \leq t + \Delta t | T_i > t\} = \frac{\Pr\{t < T_i < t + \Delta t\}}{\Pr\{T_i > t\}}.$$

The failure rate is given by the dividing this equation by the length of the time interval, Δt :

$$\frac{\Pr\{t < T_i < t + \Delta t\}}{\Delta t} \frac{1}{\Pr\{T_i > t\}} = \frac{R_{T_i}(t) - R_{T_i}(t + \Delta t)}{R_{T_i}(t)}.$$

The hazard function $h_{T_i}(t)$ for T_i is the instantaneous failure rate at time t , which is given by

$$\begin{aligned} h_{T_i}(t) &= \lim_{\Delta t \rightarrow 0} \frac{\Pr\{t < T_i < t + \Delta t\}}{\Delta t} \frac{1}{\Pr\{T_i > t\}} \\ &= \frac{f_{T_i}(t)}{R_{T_i}(t)}. \end{aligned} \quad (2)$$

\end{definition}

The lifetime of the j^{th} component is assumed to follow a parametric distribution indexed by a parameter vector $\boldsymbol{\theta}_j$. The parameter vector of the overall system is defined as

$$\boldsymbol{\theta} = (\boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_m).$$

When a random variable X is parameterized by a particular $\boldsymbol{\theta}$, we denote the reliability function by $R_X(t; \boldsymbol{\theta})$, and the same for the other distribution functions. As a special case, for the components in the series system, we subscript by their labels, e.g, the j^{th} component's pdf is denoted by $f_j(t; \boldsymbol{\theta}_j)$.

Two random variables X and Y have a joint pdf $f_{X,Y}(x, y)$. Given the joint pdf $f(x, y)$, the marginal pdf of X is given by

$$f_X(x) = \int_{\mathcal{Y}} f_{X,Y}(x, y) dy,$$

where \mathcal{Y} is the support of Y . (If Y is discrete, replace the integration with a summation over \mathcal{Y} .)

The conditional pdf of Y given $X = x$, $f_{Y|X}(y|x)$, is defined as

$$f_{Y|X}(y|x) = \frac{f_{X,Y}(x, y)}{f_X(x)}.$$

We may generalize all of the above to more than two random variables, e.g., the joint pdf of X_1, \dots, X_m is denoted by $f(x_1, \dots, x_m)$.

Next, we dive deeper into these concepts and provide mathematical derivations for the reliability function, pdf, and hazard function of the series system. We begin with the reliability function of the series system, as given by the following theorem.

Theorem 2.1. *The series system has a reliability function given by*

$$R_{T_i}(t; \boldsymbol{\theta}) = \prod_{j=1}^m R_j(t; \boldsymbol{\theta}_j). \quad (3)$$

Proof. The reliability function is defined as

$$R_{T_i}(t; \boldsymbol{\theta}) = \Pr\{T_i > t\}$$

which may be rewritten as

$$R_{T_i}(t; \boldsymbol{\theta}) = \Pr\{\min\{T_{i1}, \dots, T_{im}\} > t\}.$$

For the minimum to be larger than t , every component must be larger than t ,

$$R_{T_i}(t; \boldsymbol{\theta}) = \Pr\{T_{i1} > t, \dots, T_{im} > t\}.$$

Since the component lifetimes are independent, by the product rule the above may be rewritten as

$$R_{T_i}(t; \boldsymbol{\theta}) = \Pr\{T_{i1} > t\} \times \dots \times \Pr\{T_{im} > t\}.$$

By definition, $R_j(t; \boldsymbol{\theta}) = \Pr\{T_{ij} > t\}$. Performing this substitution obtains the result

$$R_{T_i}(t; \boldsymbol{\theta}) = \prod_{j=1}^m R_j(t; \boldsymbol{\theta}_j).$$

□

Theorem 2.1 shows that the system's overall reliability is the product of the reliabilities of its individual components. This property is inherent to series systems and will be used in the subsequent derivations.

Next, we turn our attention to the pdf of the system lifetime, described in the following theorem.

Theorem 2.2. *The series system has a pdf given by*

$$f_{T_i}(t; \boldsymbol{\theta}) = \sum_{j=1}^m f_j(t; \boldsymbol{\theta}_j) \prod_{\substack{k=1 \\ k \neq j}}^m R_k(t; \boldsymbol{\theta}_k). \quad (4)$$

Proof. By definition, the pdf may be written as

$$f_{T_i}(t; \boldsymbol{\theta}) = -\frac{d}{dt} \prod_{j=1}^m R_j(t; \boldsymbol{\theta}_j).$$

By the product rule, this may be rewritten as

$$\begin{aligned} f_{T_i}(t; \boldsymbol{\theta}) &= -\frac{d}{dt} R_1(t; \boldsymbol{\theta}_1) \prod_{j=2}^m R_j(t; \boldsymbol{\theta}_j) - R_1(t; \boldsymbol{\theta}_1) \frac{d}{dt} \prod_{j=2}^m R_j(t; \boldsymbol{\theta}_j) \\ &= f_1(t; \boldsymbol{\theta}) \prod_{j=2}^m R_j(t; \boldsymbol{\theta}_j) - R_1(t; \boldsymbol{\theta}_1) \frac{d}{dt} \prod_{j=2}^m R_j(t; \boldsymbol{\theta}_j). \end{aligned}$$

Recursively applying the product rule $m - 1$ times results in

$$f_{T_i}(t; \boldsymbol{\theta}) = \sum_{j=1}^{m-1} f_j(t; \boldsymbol{\theta}_j) \prod_{\substack{k=1 \\ k \neq j}}^m R_k(t; \boldsymbol{\theta}_k) - \prod_{j=1}^{m-1} R_j(t; \boldsymbol{\theta}_j) \frac{d}{dt} R_m(t; \boldsymbol{\theta}_m),$$

which simplifies to

$$f_{T_i}(t; \boldsymbol{\theta}) = \sum_{j=1}^m f_j(t; \boldsymbol{\theta}_j) \prod_{\substack{k=1 \\ k \neq j}}^m R_k(t; \boldsymbol{\theta}_k).$$

□

Theorem 2.2 shows the pdf of the system lifetime as a function of the pdfs and reliabilities of its components. We continue with the hazard function of the system lifetime, defined in the next theorem.

Theorem 2.3. *The series system has a hazard function given by*

$$h_{T_i}(t; \boldsymbol{\theta}) = \sum_{j=1}^m h_j(t; \boldsymbol{\theta}_j). \quad (5)$$

Proof. By Equation (2), the i^{th} series system lifetime has a hazard function defined as

$$h_{T_i}(t; \boldsymbol{\theta}) = \frac{f_{T_i}(t; \boldsymbol{\theta})}{R_{T_i}(t; \boldsymbol{\theta})}.$$

Plugging in expressions for these functions results in

$$h_{T_i}(t; \boldsymbol{\theta}) = \frac{\sum_{j=1}^m f_j(t; \boldsymbol{\theta}_j) \prod_{\substack{k=1 \\ k \neq j}}^m R_k(t; \boldsymbol{\theta}_k)}{\prod_{j=1}^m R_j(t; \boldsymbol{\theta}_j)},$$

which can be simplified to

$$h_{T_i}(t; \boldsymbol{\theta}) = \sum_{j=1}^m \frac{f_j(t; \boldsymbol{\theta}_j)}{R_j(t; \boldsymbol{\theta}_j)} = \sum_{j=1}^m h_j(t; \boldsymbol{\theta}_j).$$

□

Theorem 2.3 reveals that the system's hazard function is the sum of the hazard functions of its components. By definition, the hazard function is the ratio of the pdf to the reliability function,

$$h_{T_i}(t; \boldsymbol{\theta}) = \frac{f_{T_i}(t; \boldsymbol{\theta})}{R_{T_i}(t; \boldsymbol{\theta})},$$

and we can rearrange this to get

$$\begin{aligned} f_{T_i}(t; \boldsymbol{\theta}) &= h_{T_i}(t; \boldsymbol{\theta}) R_{T_i}(t; \boldsymbol{\theta}) \\ &= \left\{ \sum_{j=1}^m h_j(t; \boldsymbol{\theta}_j) \right\} \left\{ \prod_{j=1}^m R_j(t; \boldsymbol{\theta}_j) \right\}, \end{aligned} \quad (6)$$

which we sometimes find to be a more convenient form than Equation (4).

In this section, we derived the mathematical forms for the system's reliability function, pdf, and hazard function. Next, we build upon these concepts to derive distributions related to the component cause of failure.

2.1 Component Cause of Failure

Whenever a series system fails, precisely one of the components is the cause. We model the component cause of the series system failure as a random variable.

Definition 2.1. The component cause of failure of a series system is denoted by the random variable K_i whose support is given by $\{1, \dots, m\}$. For example, $K_i = j$ indicates that the component indexed by j failed first, i.e.,

$$T_{ij} < T_{ij'}$$

for every j' in the support of K_i except for j . Since we have series systems, K_i is unique.

The system lifetime and the component cause of failure has a joint distribution given by the following theorem.

Theorem 2.4. *The joint pdf of the component cause of failure K_i and series system lifetime T_i is given by*

$$f_{K_i, T_i}(j, t; \boldsymbol{\theta}) = h_j(t; \boldsymbol{\theta}_j) \prod_{l=1}^m R_l(t; \boldsymbol{\theta}), \quad (7)$$

where $h_j(t; \boldsymbol{\theta}_j)$ is the hazard function of the j^{th} component and $R_{T_i}(t; \boldsymbol{\theta})$ is the reliability function of the series system.

Proof. Consider a series system with 3 components. By the assumption that component lifetimes are mutually independent, the joint pdf of T_{i1}, T_{i2}, T_{i3} is given by

$$f(t_1, t_2, t_3; \boldsymbol{\theta}) = \prod_{j=1}^3 f_j(t; \boldsymbol{\theta}_j).$$

The first component is the cause of failure at time t if $K_i = 1$ and $T_i = t$, which may be rephrased as the likelihood that $T_{i1} = t$, $T_{i2} > t$, and $T_{i3} > t$. Thus,

$$\begin{aligned} f_{K_i, T_i}(j; \boldsymbol{\theta}) &= \int_t^\infty \int_t^\infty f_1(t; \boldsymbol{\theta}_1) f_2(t_2; \boldsymbol{\theta}_2) f_3(t_3; \boldsymbol{\theta}_3) dt_3 dt_2 \\ &= \int_t^\infty f_1(t; \boldsymbol{\theta}_1) f_2(t_2; \boldsymbol{\theta}_2) R_3(t; \boldsymbol{\theta}_3) dt_2 \\ &= f_1(t; \boldsymbol{\theta}_1) R_2(t; \boldsymbol{\theta}_2) R_3(t; \boldsymbol{\theta}_3). \end{aligned}$$

Since $h_1(t; \boldsymbol{\theta}_1) = f_1(t; \boldsymbol{\theta}_1)/R_1(t; \boldsymbol{\theta}_1)$,

$$f_1(t; \boldsymbol{\theta}_1) = h_1(t; \boldsymbol{\theta}_1) R_1(t; \boldsymbol{\theta}_1).$$

Making this substitution into the above expression for $f_{K_i, T_i}(j, t; \boldsymbol{\theta})$ yields

$$f_{K_i, T_i}(j, t; \boldsymbol{\theta}) = h_1(t; \boldsymbol{\theta}_1) \prod_{l=1}^m R_l(t; \boldsymbol{\theta}_l)$$

Generalizing from this completes the proof. □

The probability that the j^{th} component is the cause of failure is given by

$$\Pr\{K_i = j\} = E_{\boldsymbol{\theta}} \left[\frac{h_j(T_i; \boldsymbol{\theta}_j)}{\sum_{l=1}^m h_l(T_i; \boldsymbol{\theta}_l)} \right]. \quad (8)$$

Proof. The probability the j^{th} component is the cause of failure is given by marginalizing the joint pdf of K_i and T_i over T_i ,

$$\Pr\{K_i = j\} = \int_0^\infty f_{K_i, T_i}(j, t; \boldsymbol{\theta}) dt.$$

By Theorem 2.4, this is equivalent to

$$\begin{aligned} \Pr\{K_i = j\} &= \int_0^\infty h_j(t; \boldsymbol{\theta}_j) R_{T_i}(t; \boldsymbol{\theta}) dt \\ &= \int_0^\infty \left(\frac{h_j(t; \boldsymbol{\theta}_j)}{h_{T_i}(t; \boldsymbol{\theta})} \right) f_{T_i}(t; \boldsymbol{\theta}) dt \\ &= E_{\boldsymbol{\theta}} \left[h_j(T_i; \boldsymbol{\theta}_j) / \sum_{l=1}^m h_l(T_i; \boldsymbol{\theta}_l) \right]. \end{aligned}$$

□

2.2 System and Component Reliabilities

A common measure of reliability is mean time to failure (MTTF). The MTTF is defined as the expectation of the lifetime,

$$\text{MTTF} = E_{\boldsymbol{\theta}}\{T_i\}, \quad (9)$$

which if certain assumptions are satisfied¹ is equivalent to the integration of the reliability function over its support.

While the MTTF provides a summary measure of reliability, it is not a complete description. Depending on the failure characteristics, MTTF can be misleading. For example, a system that has a high likelihood of failing early in its life may still have a large MTTF if it is fat-tailed.²

The reliability of the components in the series system determines the reliability of the system. We denote the MTTF of the j^{th} component by MTTF_j and, according to Equation (8), the probability that the j^{th} component is the cause of failure is given by $\Pr\{K_i = j\}$. In a well-designed series system, there is no component that is the “weakest link” that either has a much shorter MTTF or a much higher probability of being the component cause of failure than any of the other components, e.g., $\Pr\{K_i = j\} \approx \Pr\{K_i = k\}$ and $\text{MTTF}_j \approx \text{MTTF}_k$ for all j and k . This just means that the components should have similar reliabilities and failure characteristics.

We use these results in the simulation study in Section 7, where we assess the sensitivity of the MLE with respect to varying the reliability of one of the Weibull components. We vary its reliability in two different ways:

1. We vary its shape parameter (keeping its scale parameter constant), which determines the failure characteristics of the component and also affects its MTTF.
2. We vary its scale parameter (keeping its shape parameter constant), which scales its MTTF while retaining the same failure characteristics.

3 Likelihood Model for Masked Data

The object of interest is the (unknown) parameter value $\boldsymbol{\theta}$. To estimate this $\boldsymbol{\theta}$, we need *data*. In our case, we call it *masked data* because we do not necessarily observe the event of interest, say a system failure, directly. We consider two types of masking: masking the system failure lifetime and masking the component cause of failure.

We generally encounter three types of system failure lifetime masking:

1. A system failure is observed at a particular point in time.
2. A system failure is observed to occur within a particular interval of time.
3. A system failure is not observed, but we know that the system survived at least until a particular point in time. This is known as *right-censoring* and can occur if, for instance, an experiment is terminated while the system is still functioning.

We generally encounter two types of component cause of failure masking:

1. The component cause of failure is observed.
2. The component cause of failure is not observed, but we know that the failed component is in some set of components. This is known as *masking* the component cause of failure.

¹ T_i is non-negative and continuous, $R_{T_i}(t; \boldsymbol{\theta})$ is a well-defined, continuous, and differential function for $t > 0$, and $\int_0^\infty R_{T_i}(t; \boldsymbol{\theta}) dt$ converges.

²A “fat-tailed” distribution refers to a probability distribution with tails that decay more slowly than those of the exponential family, such as the case with the Weibull when its shape parameter is greater than 1. This means that extreme values are more likely to occur, and the distribution is more prone to “black swan” events or rare occurrences. In the context of reliability, a fat-tailed distribution might imply a higher likelihood of unusually long lifetimes, which can skew measures like the MTTF. See, for example, Taleb (2007).

Thus, the component cause of failure masking will take the form of candidate sets. A candidate set consists of some subset of component labels that plausibly contains the label of the failed component. The sample space of candidate sets are all subsets of $\{1, \dots, m\}$, thus there are 2^m possible outcomes in the sample space.

In this paper, we limit our focus to observing *right censored* lifetimes and exact lifetimes but with masked component cause of failures. We consider a sample of n i.i.d. series systems, each of which is put into operation at some time and and observed until either it fails or is right-censored. We denote the right-censoring time of the i^{th} system by τ_i . We do not directly observe the system lifetime, T_i , but rather, we observe the right-censored lifetime, S_i , which is given by

$$S_i = \min\{\tau_i, T_i\}, \quad (10)$$

We also observe a right-censoring indicator, δ_i , which is given by

$$\delta_i = 1_{T_i < \tau_i} \quad (11)$$

where $1_{\text{condition}}$ is an indicator function that outputs 1 if *condition* is true and 0 otherwise. Here, $\delta_i = 1$ indicates the event of interest, a system failure, was observed.

If a system failure lifetime is observed, then we also observe a candidate set that contains the component cause of failure. We denote the candidate set for the i^{th} system by \mathcal{C}_i , which is a subset of $\{1, \dots, m\}$. Since the data generating process for candidate sets may be subject to chance variations, it as a random set.

Consider we have an independent and identically distributed (i.i.d.) random sample of masked data, $D = \{D_1, \dots, D_n\}$, where each D_i contains the following:

- S_i , the system lifetime of the i^{th} system.
- δ_i , the right-censoring indicator of the i^{th} system.
- \mathcal{C}_i , the set of candidate component causes of failure for the i^{th} system.

The masked data generation process is illustrated by Figure 1.




Figure 1: This figure showcases a dependency graph of the generative model for $D_i = (S_i, \delta_i, \mathcal{C}_i)$. The elements in green are observed in the sample, while the elements in red are unobserved (latent). We see that \mathcal{C}_i is related to both the unobserved component lifetimes T_{i1}, \dots, T_{im} and other unknown and unobserved covariates, like ambient temperature or the particular diagnostician who generated the candidate set. These two complications for \mathcal{C}_i are why seek a way to construct a reduced likelihood function in later sections that is not a function of the distribution of \mathcal{C}_i .

An example of masked data D for exact, right-censored system failure times with candidate sets that mask the component cause of failure can be seen in Table 1 for a series system with $m = 3$ components.

Table 1: Right-censored lifetime data with masked component cause of failure.

System	Right-censoring time (S_i)	Right censoring indicator (δ_i)	Candidate set (\mathcal{C}_i)
1	4.3	1	$\{1, 2\}$
2	1.3	1	$\{2\}$
3	5.4	0	\emptyset
4	2.6	1	$\{2, 3\}$
5	3.7	1	$\{1, 2, 3\}$
6	10	0	\emptyset

In our model, we assume the data is governed by a pdf, which is determined by a specific parameter, represented as θ within the parameter space Ω . The joint pdf of the data D can be represented as follows:

$$f(D; \theta) = \prod_{i=1}^n f(s_i, \delta_i, c_i; \theta),$$

where s_i is the observed system lifetime of the i^{th} system, δ_i is the observed right-censoring indicator of the i^{th} system, and c_i is the observed candidate set of the i^{th} system.

This joint pdf tells us how likely we are to observe the particular data, D , given the parameter θ . When we keep the data constant and allow the parameter θ to vary, we obtain what is called the likelihood function L , defined as

$$L(\theta) = \prod_{i=1}^n L_i(\theta)$$

where

$$L_i(\theta) = f(s_i, \delta_i, c_i; \theta)$$

is the likelihood contribution of the i^{th} system. In other words, the likelihood function quantifies how likely different parameter values θ are, given the observed data.

For each type of data, right-censored data and masked component cause of failure data, we will derive the *likelihood contribution* L_i , which refers to the part of the likelihood function that this particular piece of data contributes to.

We present the following theorem for the likelihood contribution model.

Theorem 3.1. *The likelihood contribution of the i -th system is given by*

$$L_i(\theta) = R_{T_i}(s_i; \theta) \left(\beta_i \sum_{j \in c_i} h_j(s_i; \theta_j) \right)^{\delta_i} \quad (12)$$

where $R_{T_i}(s_i; \theta) = \prod_{j=1}^m R_j(s_i; \theta_j)$ is the reliability function of the series system evaluated at s_i , $\delta_i = 0$ indicates the i^{th} system is right-censored at time s_i , and $\delta_i = 1$ indicates the i^{th} system is observed to have failed at time s_i with a component cause of failure is masked by the candidate set c_i .

In the follow subsections, we prove this result for each type of masked data, right-censored system lifetime data ($\delta_i = 0$) and masking of the component cause of failure ($\delta_i = 1$).

3.1 Masked Component Cause of Failure

Suppose a diagnostician is unable to identify the precise component cause of the failure, e.g., due to cost considerations he or she replaced multiple components at once, successfully repairing the system but failing to precisely identify the failed component. In this case, the cause of failure is said to be *masked*.

The unobserved component lifetimes may have many covariates, like ambient operating temperature, but the only covariate we observe in our masked data model are the system's lifetime and additional masked data in the form of a candidate set that is somehow correlated with the unobserved component lifetimes.

The key goal of our analysis is to estimate the parameters, θ , which maximize the likelihood of the observed data, and to estimate the precision and accuracy of this estimate using the Bootstrap method.

To achieve this, we first need to assess the joint distribution of the system's continuous lifetime, T_i , and the discrete candidate set, \mathcal{C}_i , which can be written as

$$f_{T_i, \mathcal{C}_i}(t_i, c_i; \theta) = f_{T_i}(t_i; \theta) \Pr_{\theta}\{\mathcal{C}_i = c_i | T_i = t_i\},$$

where $f_{T_i}(t_i; \theta)$ is the pdf of T_i and $\Pr_{\theta}\{\mathcal{C}_i = c_i | T_i = t_i\}$ is the conditional pmf of \mathcal{C}_i given $T_i = t_i$.

We assume the pdf $f_{T_i}(t_i; \theta)$ is known, but we do not have knowledge of $\Pr_{\theta}\{\mathcal{C}_i = c_i | T_i = t_i\}$, i.e., the data generating process for candidate sets is unknown.

However, it is critical that the masked data, \mathcal{C}_i , is correlated with the i^{th} system. This way, the conditional distribution of \mathcal{C}_i given $T_i = t_i$ may provide information about θ , despite our Statistical interest being primarily in the series system rather than the candidate sets.

To make this problem tractable, we assume a set of conditions that make it unnecessary to estimate the generative processes for candidate sets. The most important way in which \mathcal{C}_i is correlated with the i^{th} system is given by assuming the following condition.

Condition 1. The candidate set \mathcal{C}_i contains the index of the the failed component, i.e.,

$$\Pr_{\theta}\{K_i \in \mathcal{C}_i\} = 1$$

where K_i is the random variable for the failed component index of the i^{th} system.

Assuming Condition 1, \mathcal{C}_i must contain the index of the failed component, but we can say little else about what other component indices may appear in \mathcal{C}_i .

In order to derive the joint distribution of \mathcal{C}_i and T_i assuming Condition 1, we take the following approach. We notice that \mathcal{C}_i and K_i are statistically dependent. We denote the conditional pmf of \mathcal{C}_i given $T_i = t_i$ and $K_i = j$ as

$$\Pr_{\theta}\{\mathcal{C}_i = c_i | T_i = t_i, K_i = j\}.$$

Even though K_i is not observable in our masked data model, we can still consider the joint distribution of T_i , K_i , and \mathcal{C}_i . By Theorem 2.4, the joint pdf of T_i and K_i is given by

$$f_{T_i, K_i}(t_i, j; \theta) = h_j(t_i; \theta_j) R_{T_i}(t_i; \theta),$$

where $h_j(t_i; \theta_j)$ is the hazard function for the j^{th} component and $R_{T_i}(t_i; \theta)$ is the reliability function of the system. Thus, the joint pdf of T_i , K_i , and \mathcal{C}_i may be written as

$$\begin{aligned} f_{T_i, K_i, \mathcal{C}_i}(t_i, j, c_i; \theta) &= f_{T_i, K_i}(t_i, j; \theta) \Pr_{\theta}\{\mathcal{C}_i = c_i | T_i = t_i, K_i = j\} \\ &= h_j(t_i; \theta_j) R_{T_i}(t_i; \theta) \Pr_{\theta}\{\mathcal{C}_i = c_i | T_i = t_i, K_i = j\}. \end{aligned} \quad (13)$$

We are going to need the joint pdf of T_i and \mathcal{C}_i , which may be obtained by summing over the support $\{1, \dots, m\}$ of K_i in Equation (13),

$$f_{T_i, \mathcal{C}_i}(t_i, c_i; \theta) = R_{T_i}(t_i; \theta) \sum_{j=1}^m \left\{ h_j(t_i; \theta_j) \Pr_{\theta}\{\mathcal{C}_i = c_i | T_i = t_i, K_i = j\} \right\}.$$

By Condition 1, $\Pr_{\theta}\{\mathcal{C}_i = c_i | T_i = t_i, K_i = j\} = 0$ when $K_i = j$ and $j \notin c_i$, and so we may rewrite the joint pdf of T_i and \mathcal{C}_i as

$$f_{T_i, \mathcal{C}_i}(t_i, c_i; \theta) = R_{T_i}(t_i; \theta) \sum_{j \in c_i} \left\{ h_j(t_i; \theta_j) \Pr_{\theta}\{\mathcal{C}_i = c_i | T_i = t_i, K_i = j\} \right\}. \quad (14)$$

When we try to find an MLE of θ (see Section 4), we solve the simultaneous equations of the MLE and choose a solution $\hat{\theta}$ that is a maximum for the likelihood function. When we do this, we find that $\hat{\theta}$ depends on the unknown conditional pmf $\Pr_{\theta}\{\mathcal{C}_i = c_i | T_i = t_i, K_i = j\}$. So, we are motivated to seek out more conditions (that approximately hold in realistic situations) whose MLEs are independent of the pmf $\Pr_{\theta}\{\mathcal{C}_i = c_i | T_i = t_i, K_i = j\}$.

Condition 2. Any of the components in the candidate set has an equal probability of being the cause of failure. That is, for a fixed $j \in c_i$,

$$\Pr_{\theta}\{\mathcal{C}_i = c_i | T_i = t_i, K_i = j'\} = \Pr_{\theta}\{\mathcal{C}_i = c_i | T_i = t_i, K_i = j\}$$

for all $j' \in c_i$.

According to (Guess et al., 1991), in many industrial problems, masking generally occurred due to time constraints and the expense of failure analysis. In this setting, Condition 2 generally holds.

Assuming Conditions 1 and 2, $\Pr_{\theta}\{\mathcal{C}_i = c_i | T_i = t_i, K_i = j\}$ may be factored out of the summation in Equation (14), and thus the joint pdf of T_i and \mathcal{C}_i may be rewritten as

$$f_{T_i, \mathcal{C}_i}(t_i, c_i; \theta) = \Pr_{\theta}\{\mathcal{C}_i = c_i | T_i = t_i, K_i = j'\} R_{T_i}(t_i; \theta) \sum_{j \in c_i} h_j(t_i; \theta_j)$$

where $j' \in c_i$.

If $\Pr_{\theta}\{\mathcal{C}_i = c_i | T_i = t_i, K_i = j'\}$ is a function of θ , the MLEs are still dependent on the unknown $\Pr_{\theta}\{\mathcal{C}_i = c_i | T_i = t_i, K_i = j'\}$. This is a more tractable problem, but we are primarily interested in the situation where we do not need to know (nor estimate) $\Pr_{\theta}\{\mathcal{C}_i = c_i | T_i = t_i, K_i = j'\}$ to find an MLE of θ . The last condition we assume achieves this result.

Condition 3. The masking probabilities conditioned on failure time T_i and component cause of failure K_i are not functions of θ . In this case, the conditional probability of \mathcal{C}_i given $T_i = t_i$ and $K_i = j'$ is denoted by

$$\beta_i = \Pr\{\mathcal{C}_i = c_i | T_i = t_i, K_i = j'\}$$

where β_i is not a function of θ .

When Conditions 1, 2, and 3 are satisfied, the joint pdf of T_i and \mathcal{C}_i is given by

$$f_{T_i, \mathcal{C}_i}(t_i, c_i; \theta) = \beta_i R_{T_i}(t_i; \theta) \sum_{j \in c_i} h_j(t_i; \theta_j).$$

When we fix the sample and allow θ to vary, we obtain the contribution to the likelihood L from the i^{th} observation when the system lifetime is exactly known (i.e., $\delta_i = 1$) but the component cause of failure is masked by a candidate set c_i :

$$L_i(\theta) = R_{T_i}(t_i; \theta) \sum_{j \in c_i} h_j(t_i; \theta_j). \quad (15)$$

To summarize this result, assuming Conditions 1, 2, and 3, if we observe an exact system failure time for the i -th system ($\delta_i = 1$), but the component that failed is masked by a candidate set c_i , then its likelihood contribution is given by Equation (15).

3.2 Right-Censored Data

As described in Section 3, we observe realizations of $(S_i, \delta_i, \mathcal{C}_i)$ where $S_i = \min\{T_i, \tau_i\}$ is the right-censored system lifetime, $\delta_i = 1_{\{T_i < \tau_i\}}$ is the right-censoring indicator, and \mathcal{C}_i is the candidate set.

In the previous section, we discussed the likelihood contribution from an observation of a masked component cause of failure, i.e., $\delta_i = 1$. We now derive the likelihood contribution of a *right-censored* observation ($\delta_i = 0$) in our masked data model.

Theorem 3.2. *The likelihood contribution of a right-censored observation ($\delta_i = 0$) is given by*

$$L_i(\theta) = R_{T_i}(s_i; \theta). \quad (16)$$

Proof. When right-censoring occurs, then $S_i = \tau_i$, and we only know that $T_i > \tau_i$, and so we integrate over all possible values that it may have obtained,

$$L_i(\theta) = \Pr_{\theta}\{T_i > s_i\}.$$

By definition, this is just the survival or reliability function of the series system evaluated at s_i ,

$$L_i(\theta) = R_{T_i}(s_i; \theta).$$

□

When we combine the two likelihood contributions, we obtain the likelihood contribution for the i^{th} system shown in Theorem 3.1,

$$L_i(\boldsymbol{\theta}) = \begin{cases} R_{T_i}(s_i; \boldsymbol{\theta}) & \text{if } \delta_i = 0 \\ \beta_i R_{T_i}(s_i; \boldsymbol{\theta}) \sum_{j \in c_i} h_j(s_i; \boldsymbol{\theta}_j) & \text{if } \delta_i = 1. \end{cases}$$

We use this result in Section 4 to derive the maximum likelihood estimator (MLE) of $\boldsymbol{\theta}$.

3.3 Identifiability and Convergence Issues

In our likelihood model, masking and right-censoring can lead to issues related to identifiability and flat likelihood regions. Identifiability refers to the unique mapping of the model parameters to the likelihood function, and lack of identifiability can lead to multiple sets of parameters that explain the data equally well, making inference about the true parameters challenging (Lehmann and Casella, 1998), while flat likelihood regions can complicate convergence (Wu, 1983).

In our simulation study, we address these challenges in a pragmatic way. Specifically, failure to converge to a solution within a maximum of 125 iterations is interpreted as evidence of the aforementioned issues, leading to the discarding of the sample, with the process then repeated with a new synthetic sample. Note, however, that in Section 5 where we discuss the bias-corrected and accelerated (BCa) bootstrap method for constructing confidence intervals, we do not discard any resamples.

This strategy helps ensure the robustness of the results, while acknowledging the inherent complexities of likelihood-based estimation in models characterized by masking and right-censoring.

4 Maximum Likelihood Estimation

In our analysis, we use maximum likelihood estimation (MLE) to estimate the series system parameter $\boldsymbol{\theta}$ from the masked data (Bain and Engelhardt, 1992; Casella and Berger, 2002). The MLE finds parameter values that maximize the likelihood of the observed data under the assumed model. A maximum likelihood estimate, $\hat{\boldsymbol{\theta}}$, is a solution of

$$L(\hat{\boldsymbol{\theta}}) = \max_{\boldsymbol{\theta} \in \Omega} L(\boldsymbol{\theta}), \quad (17)$$

where $L(\boldsymbol{\theta})$ is the likelihood function of the observed data. For computational efficiency and analytical simplicity, we work with the log-likelihood function, denoted as $\ell(\boldsymbol{\theta})$, instead of the likelihood function (Casella and Berger, 2002).

Theorem 4.1. *The log-likelihood function, $\ell(\boldsymbol{\theta})$, for our masked data model is the sum of the log-likelihoods for each observation,*

$$\ell(\boldsymbol{\theta}) = \sum_{i=1}^n \ell_i(\boldsymbol{\theta}), \quad (18)$$

where $\ell_i(\boldsymbol{\theta})$ is the log-likelihood contribution for the i^{th} observation:

$$\ell_i(\boldsymbol{\theta}) = \sum_{j=1}^m \log R_j(s_i; \boldsymbol{\theta}_j) + \delta_i \log \left(\sum_{j \in c_i} h_j(s_i; \boldsymbol{\theta}_j) \right). \quad (19)$$

Proof. The log-likelihood function is the logarithm of the likelihood function,

$$\ell(\boldsymbol{\theta}) = \log L(\boldsymbol{\theta}) = \log \prod_{i=1}^n L_i(\boldsymbol{\theta}) = \sum_{i=1}^n \log L_i(\boldsymbol{\theta}).$$

Substituting $L_i(\boldsymbol{\theta})$ from Equation (12), we consider these two cases of δ_i separately to obtain the result in Theorem 4.1.

Case 1: If the i -th system is right-censored ($\delta_i = 0$),

$$\ell_i(\boldsymbol{\theta}) = \log R_{T_i}(s_i; \boldsymbol{\theta}) = \sum_{j=1}^m \log R_j(s_i; \boldsymbol{\theta}_j).$$

Case 2: If the i -th system's component cause of failure is masked but the failure time is known ($\delta_i = 1$),

$$\ell_i(\boldsymbol{\theta}) = \log R_{T_i}(s_i; \boldsymbol{\theta}) + \log \beta_i + \log \left(\sum_{j \in c_i} h_j(s_i; \boldsymbol{\theta}_j) \right).$$

We replace $R_{T_i}(s_i; \boldsymbol{\theta})$ with its component-wise definition and by Condition 3, we may discard³ the $\log \beta_i$ term since it does not depend on $\boldsymbol{\theta}$, giving us the result

$$\ell_i(\boldsymbol{\theta}) = \sum_{j=1}^m \log R_j(s_i; \boldsymbol{\theta}_j) + \log \left(\sum_{j \in c_i} h_j(s_i; \boldsymbol{\theta}_j) \right).$$

Combining these two cases gives us the result in Theorem 4.1. □

The MLE, $\hat{\boldsymbol{\theta}}$, is often found by solving a system of equations derived from setting the derivative of the log-likelihood function to zero, i.e.,

$$\frac{\partial}{\partial \theta_j} \ell(\boldsymbol{\theta}) = 0, \tag{20}$$

for each component θ_j of the parameter $\boldsymbol{\theta}$ (Bain and Engelhardt, 1992). When there's no closed-form solution, we resort to numerical methods like the Newton-Raphson method.

Assuming some regularity conditions, such as the likelihood function being identifiable, the MLE has many desirable asymptotic properties that underpin statistical inference, namely that it is an asymptotically unbiased estimator of the parameter $\boldsymbol{\theta}$ and it is normally distributed with a variance given by the inverse of the Fisher Information Matrix (FIM) (Casella and Berger, 2002). However, for smaller samples, these asymptotic properties may not yield accurate approximations. We propose to use the bootstrap method to offer an empirical approach for estimating the sampling distribution of the MLE, in particular for computing confidence intervals.

5 Bias-Corrected and Accelerated Bootstrap Confidence Intervals

We utilize the non-parametric bootstrap to approximate the sampling distribution of the MLE. In the non-parametric bootstrap, we resample from the observed data with replacement to generate a bootstrap sample. The MLE is then computed for the bootstrap sample. This process is repeated B times, giving us B bootstrap replicates of the MLE. The sampling distribution of the MLE is then approximated by the empirical distribution of the bootstrap replicates of the MLE.

The method we use to generate confidence intervals is known as Bias-Corrected and Accelerated Bootstrap Confidence Intervals (BCa), which applies two corrections to the standard bootstrap method:

- **Bias correction:** This adjusts for bias in the bootstrap distribution itself. This bias is measured as the difference between the mean of the bootstrap distribution and the observed statistic. It works by transforming the percentiles of the bootstrap distribution to correct for these issues.

This may be a useful transformation in our case since we are dealing with small samples and we have two potential sources of bias: right-censoring and masking component cause of failure. They seem to have opposing effects on the MLE, but the relationship is difficult to quantify.

³Adding or subtracting a function by a constant does not change where it obtains a maximum, so we are free to discard such terms from the log-likelihood function.

- **Acceleration:** This adjusts for the rate of change of the statistic as a function of the true, unknown parameter. This correction is important when the shape of the statistic’s distribution changes with the true parameter.

Since we have a number of different shape parameters, k_1, \dots, k_m , we may expect the shape of the distribution of the MLE to change as a function of the true parameter, making this correction potentially useful.

Since we are primarily interested in generating confidence intervals for small samples for a potentially biased MLE, the BCa method may be a good choice for our analysis. For more details on BCa, see Efron (1987).

In our simulation study, we will assess the performance of the bootstrapped BCa confidence intervals by computing the coverage probability of the confidence intervals. A well-calibrated 95% confidence interval contains the true value around 95% of the time. If the confidence interval is too narrow, it will have a coverage probability less than 95%, which conveys a sort of false confidence in the precision of the MLE. If the confidence interval is too wide, it will have a coverage probability greater than 95%, which conveys a lack of confidence in the precision of the MLE. We want confidence intervals to be as narrow as possible while still having a coverage probability close to the nominal level, 95%.

5.1 Issues with Resampling from the Observed Data

While the bootstrap method provides a robust and flexible tool for statistical estimation, its effectiveness can be influenced by several factors (Efron and Tibshirani, 1994).

Firstly, instances of non-convergence in our bootstrap samples were observed. Such cases can occur when the estimation method, like the MLE used in our analysis, fails to converge due to the specifics of the resampled data (Casella and Berger, 2002). This issue can potentially introduce bias or reduce the effective sample size of our bootstrap distribution.

Secondly, the bootstrap’s accuracy can be compromised with small sample sizes, as the method relies on the law of large numbers to approximate the true sampling distribution. For small datasets, the bootstrap samples might not adequately represent the true variability in the data, leading to inaccurate results (Efron and Tibshirani, 1994).

Thirdly, our data involves right censoring and a masking of the component cause of failure when a system failure is observed. These aspects can cause certain data points or trends to be underrepresented or not represented at all in our data, introducing bias in the bootstrap distribution (Klein and Moeschberger, 2005).

Despite these challenges, we found the bootstrap method useful in approximating the sampling distribution of the MLE, taking care in interpreting the results, particularly as it relates to coverage probabilities.

6 Series System with Weibull Components

The Weibull distribution, introduced by Waloddi Weibull in 1937, has been instrumental in reliability analysis due to its ability to model a wide range of failure behaviors. Reflecting on its utility, Weibull modestly noted that it “[...] may sometimes render good service.” (Abernethy, 2006). In the context of our study, we utilize the Weibull to model a system as originating from Weibull components in a series configuration, producing a specific form of the likelihood model described in Section 3, which deals with challenges such as right censoring and masked component cause of failure.

The j^{th} component of the i^{th} has a lifetime distribution given by

$$T_{ij} \sim \text{Weibull}(k_j, \lambda_j) \quad \text{for } i = 1, \dots, n \text{ and } j = 1, \dots, m,$$

where $\lambda_j > 0$ is the scale parameter and $k_j > 0$ is the shape parameter. The j^{th} component has a reliability

function, pdf, and hazard function given respectively by

$$R_j(t; \lambda_j, k_j) = \exp\left\{-\left(\frac{t}{\lambda_j}\right)^{k_j}\right\}, \quad (21)$$

$$f_j(t; \lambda_j, k_j) = \frac{k_j}{\lambda_j} \left(\frac{t}{\lambda_j}\right)^{k_j-1} \exp\left\{-\left(\frac{t}{\lambda_j}\right)^{k_j}\right\}, \quad (22)$$

$$h_j(t; \lambda_j, k_j) = \frac{k_j}{\lambda_j} \left(\frac{t}{\lambda_j}\right)^{k_j-1}. \quad (23)$$

The shape parameter of the Weibull distribution is of particular importance:

- $k_j < 1$ indicates infant mortality. An example of how this might arise is a result of defective components being weeded out early, and the remaining components surviving for a much longer time.
- $k_j = 1$ indicates random failures (independent of age). An example of how this might arise is a result of random shocks to the system, but otherwise the system is age-independent.⁴
- $k_j > 1$ indicates wear-out failures. An example of how this might arise is a result of components wearing as they age

The lifetime of the series system composed of m Weibull components has a reliability and hazard functions given by the following theorem.

Theorem 6.1. *The lifetime of a series system composed of m Weibull components has a reliability function, hazard function, and pdf respectively given by*

$$R_{T_i}(t; \boldsymbol{\theta}) = \exp\left\{-\sum_{j=1}^m \left(\frac{t}{\lambda_j}\right)^{k_j}\right\}, \quad (24)$$

$$h_{T_i}(t; \boldsymbol{\theta}) = \sum_{j=1}^m \frac{k_j}{\lambda_j} \left(\frac{t}{\lambda_j}\right)^{k_j-1}, \quad (25)$$

$$f_{T_i}(t; \boldsymbol{\theta}) = \left\{ \sum_{j=1}^m \frac{k_j}{\lambda_j} \left(\frac{t}{\lambda_j}\right)^{k_j-1} \right\} \exp\left\{-\sum_{j=1}^m \left(\frac{t}{\lambda_j}\right)^{k_j}\right\}. \quad (26)$$

Proof. The proof for the reliability function follows from Theorem 2.1,

$$R_{T_i}(t; \boldsymbol{\theta}) = \prod_{j=1}^m R_j(t; \lambda_j, k_j).$$

Plugging in the Weibull component reliability functions obtains the result

$$R_{T_i}(t; \boldsymbol{\theta}) = \prod_{j=1}^m \exp\left\{-\left(\frac{t}{\lambda_j}\right)^{k_j}\right\} = \exp\left\{-\sum_{j=1}^m \left(\frac{t}{\lambda_j}\right)^{k_j}\right\}.$$

The proof for the hazard function follows from Theorem 2.3,

$$h_{T_i}(t; \boldsymbol{\theta}) = \sum_{j=1}^m h_j(t; \boldsymbol{\theta}_j) = \sum_{j=1}^m \frac{k_j}{\lambda_j} \left(\frac{t}{\lambda_j}\right)^{k_j-1}$$

The proof for the pdf follows from Theorem 2.2. By definition,

$$f_{T_i}(t; \boldsymbol{\theta}) = h_{T_i}(t; \boldsymbol{\theta}) R_{T_i}(t; \boldsymbol{\theta}).$$

Plugging in the failure rate and reliability functions given respectively by Equations (24) and (25) completes the proof. \square

⁴The exponential distribution is a special case of the Weibull distribution when $k_j = 1$.

6.1 Reliability

In Section 2.2, we discussed the concept of reliability. In the case of Weibull components, the MTTF of the j^{th} component is given by

$$\text{MTTF}_j = \lambda_j \Gamma\left(1 + \frac{1}{k_j}\right), \quad (27)$$

where Γ is the gamma function.

We mentioned that the MTTF can sometimes be a poor measure of reliability, e.g., the MTTF and the probability of failing early can be large. The Weibull is a good example of this phenomenon. If $k > 1$, the Weibull is a fat-tailed distribution, and it can exhibit both a large MTTF and a high probability of failing early.

Components may have similar MTTFs, but some components may be more likely to fail early and others may be more likely to fail late, depending upon their failure characteristics (shape parameters), and so the probability of component failure given by Equation (8) is a useful measure of component reliability compared to the other components in the system.

In a well-designed series system, the component failure characteristics are similar: they have a similar MTTF and a similar probability of being the component cause of failure, i.e., they have similar shapes and scales, so that system failures are not dominated by some subset of components.

6.2 Likelihood Model

In Section 3, we discussed two separate kinds of likelihood contributions, masked component cause of failure data (with exact system failure times) and right-censored data. The likelihood contribution of the i^{th} system is given by the following theorem.

Theorem 6.2. *Let δ_i be an indicator variable that is 1 if the i^{th} system fails and 0 (right-censored) otherwise. Then the likelihood contribution of the i^{th} system is given by*

$$L_i(\boldsymbol{\theta}) = \begin{cases} \exp\left\{-\sum_{j=1}^m \left(\frac{t_i}{\lambda_j}\right)^{k_j}\right\} \beta_i \sum_{j \in c_i} \frac{k_j}{\lambda_j} \left(\frac{t_i}{\lambda_j}\right)^{k_j-1} & \text{if } \delta_i = 1, \\ \exp\left\{-\sum_{j=1}^m \left(\frac{t_i}{\lambda_j}\right)^{k_j}\right\} & \text{if } \delta_i = 0. \end{cases} \quad (28)$$

Proof. By Theorem 3.1, the likelihood contribution of the i -th system is given by

$$L_i(\boldsymbol{\theta}) = \begin{cases} R_{T_i}(s_i; \boldsymbol{\theta}) & \text{if } \delta_i = 0 \\ \beta_i R_{T_i}(s_i; \boldsymbol{\theta}) \sum_{j \in c_i} h_j(s_i; \boldsymbol{\theta}_j) & \text{if } \delta_i = 1. \end{cases}$$

By Equation (24), the system reliability function R_{T_i} is given by

$$R_{T_i}(t_i; \boldsymbol{\theta}) = \exp\left\{-\sum_{j=1}^m \left(\frac{t_i}{\lambda_j}\right)^{k_j}\right\}.$$

and by Equation (23), the Weibull component hazard function h_j is given by

$$h_j(t_i; \boldsymbol{\theta}_j) = \frac{k_j}{\lambda_j} \left(\frac{t_i}{\lambda_j}\right)^{k_j-1}.$$

Plugging these into the likelihood contribution function obtains the result. \square

Taking the log of the likelihood contribution function obtains the following result.

Corollary 6.1. *The log-likelihood contribution of the i -th system is given by*

$$\ell_i(\boldsymbol{\theta}) = -\sum_{j=1}^m \left(\frac{t_i}{\lambda_j}\right)^{k_j} + \delta_i \log\left(\sum_{j \in c_i} \frac{k_j}{\lambda_j} \left(\frac{t_i}{\lambda_j}\right)^{k_j-1}\right) \quad (29)$$

where we drop any terms that do not depend on $\boldsymbol{\theta}$ since they do not affect the MLE.

We find an MLE by solving (20), i.e., a point $\hat{\boldsymbol{\theta}} = (\hat{k}_1, \hat{\lambda}_1, \dots, \hat{k}_m, \hat{\lambda}_m)$ satisfying $\nabla_{\boldsymbol{\theta}} \ell(\hat{\boldsymbol{\theta}}) = \mathbf{0}$, where $\nabla_{\boldsymbol{\theta}}$ is the gradient of the log-likelihood function (score) with respect to $\boldsymbol{\theta}$.

To solve this system of equations, we use the Newton-Raphson method, which requires the score and the Hessian of the log-likelihood function. We analytically derive the score since it is useful to have for the Newton-Raphson method, but we do not do the same for the Hessian of the log-likelihood for the following reasons:

1. The gradient is relatively easy to derive, and it is useful to have for computing gradients efficiently and accurately, which will be useful for numerically approximating the Hessian.
2. The Hessian is tedious and error prone to derive, and Newton-like methods often do not require the Hessian to be explicitly computed.

The following theorem derives the score function.

Theorem 6.3. *The score function of the log-likelihood contribution of the i -th Weibull series system is given by*

$$\nabla \ell_i(\boldsymbol{\theta}) = \left(\frac{\partial \ell_i(\boldsymbol{\theta})}{\partial k_1}, \frac{\partial \ell_i(\boldsymbol{\theta})}{\partial \lambda_1}, \dots, \frac{\partial \ell_i(\boldsymbol{\theta})}{\partial k_m}, \frac{\partial \ell_i(\boldsymbol{\theta})}{\partial \lambda_m} \right)', \quad (30)$$

where

$$\frac{\partial \ell_i(\boldsymbol{\theta})}{\partial k_r} = - \left(\frac{t_i}{\lambda_r} \right)^{k_r} \log \left(\frac{t_i}{\lambda_r} \right) + \frac{\frac{1}{t_i} \left(\frac{t_i}{\lambda_r} \right)^{k_r} (1 + k_r \log \left(\frac{t_i}{\lambda_r} \right))}{\sum_{j \in c_i} \frac{k_j}{\lambda_j} \left(\frac{t_i}{\lambda_j} \right)^{k_j - 1}} 1_{\delta_i = 1 \wedge r \in c_i} \quad (31)$$

and

$$\frac{\partial \ell_i(\boldsymbol{\theta})}{\partial \lambda_r} = \frac{k_r}{\lambda_r} \left(\frac{t_i}{\lambda_r} \right)^{k_r} - \frac{\left(\frac{k_r}{\lambda_r} \right)^2 \left(\frac{t_i}{\lambda_r} \right)^{k_r - 1}}{\sum_{j \in c_i} \frac{k_j}{\lambda_j} \left(\frac{t_i}{\lambda_j} \right)^{k_j - 1}} 1_{\delta_i = 1 \wedge r \in c_i} \quad (32)$$

The result follows from taking the partial derivatives of the log-likelihood contribution of the i -th system given by Equation (28). It is a tedious calculation so the proof has been omitted, but the result has been verified by using a very precise numerical approximation of the gradient.

By the linearity of differentiation, the gradient of a sum of functions is the sum of their gradients, and so the score function conditioned on the entire sample is given by

$$\nabla \ell(\boldsymbol{\theta}) = \sum_{i=1}^n \nabla \ell_i(\boldsymbol{\theta}). \quad (33)$$

6.3 Weibull Series System: Homogeneous Shape Parameters

A series system composed of Weibull components is not generally Weibull unless the shape parameters of the components are homogeneous.

Theorem 6.4. *If the shape parameters of the components are homogenous, then the series system is Weibull with a shape parameter k given by the identical shape parameters of the components and a scale parameter λ given by*

$$\lambda = \left(\sum_{j=1}^m \lambda_j^{-k} \right)^{-1/k}, \quad (34)$$

where λ_j is the scale parameter of the j^{th} component.

Proof. Given m Weibull lifetimes T_{i1}, \dots, T_{im} with the same shape parameter k and scale parameters $\lambda_1, \dots, \lambda_m$, the reliability function of the series system is

$$R_{T_i}(t; \boldsymbol{\theta}) = \exp \left\{ - \sum_{j=1}^m \left(\frac{t}{\lambda_j} \right)^k \right\}.$$

To make this a Weibull system, we need to find a single scale parameter λ such that

$$R_{T_i}(t; \boldsymbol{\theta}) = \exp \left\{ - \left(\frac{t}{\lambda} \right)^k \right\},$$

which has the solution

$$\lambda = \frac{1}{\left(\frac{1}{\lambda_1^k} + \dots + \frac{1}{\lambda_m^k} \right)^{\frac{1}{k}}}.$$

□

Theorem 6.5. *If a series system has Weibull components with homogeneous shape parameters, the component cause of failure is conditionally independent of the system failure time:*

$$\Pr\{K_i = j | T_i = t_i\} = \Pr\{K_i = j\} = \frac{\lambda_j^{-k}}{\sum_{l=1}^m \lambda_l^{-k}}.$$

Proof. The conditional probability of the j^{th} component being the cause of failure given the system failure time is given by

$$\begin{aligned} \Pr\{K_i = j | T_i = t\} &= \frac{f_{K_i, T_i}(j, t; \boldsymbol{\theta})}{f_{T_i}(t; \boldsymbol{\theta})} = \frac{h_j(t; k, \lambda_j) R_{T_i}(t; \boldsymbol{\theta})}{h_{T_i}(t; \boldsymbol{\theta}_j) R_{T_i}(t; \boldsymbol{\theta})} \\ &= \frac{h_j(t; k, \lambda_j)}{\sum_{l=1}^m h_l(t; k, \lambda_l)} = \frac{\frac{k}{\lambda_j} \left(\frac{t}{\lambda_j} \right)^{k-1}}{\sum_{l=1}^m \frac{k}{\lambda_l} \left(\frac{t}{\lambda_l} \right)^{k-1}} = \frac{\left(\frac{1}{\lambda_j} \right)^k}{\sum_{l=1}^m \left(\frac{1}{\lambda_l} \right)^k}. \end{aligned}$$

□

If we have prior knowledge that the shape parameters are sufficiently homogenous, it may be useful to simplify the likelihood model by assuming the shape parameters are identical, simplifying the series system to Weibull, facilitating analysis and interpretation. According to the bias-variance trade-off, we expect the MLE to be more biased but have lower sampling variance.

We denote the full model log-likelihood function by ℓ_F and the reduced model log-likelihood by ℓ_R . The reduced model is obtained by setting the shape parameter of each component to be the same, i.e., $k_1 = \dots = k_m = k$. Thus, the reduced model log-likelihood function is given by

$$\ell_R(k, \lambda_1, \lambda_2, \dots, \lambda_m) = \ell_F(k, \lambda_1, k, \lambda_2, \dots, k, \lambda_m),$$

The same may be done for the score and hessian of the log-likelihood functions.

7 Simulation Study: Full Model

In this section, we conduct a simulation study to assess the performance of the MLE for the full likelihood model defined in Section 6. In this simulation study, we assess the sensitivity of the MLE to various simulation scenarios. In particular, we assess two important properties of the MLE with respect to a scenario:

1. Accuracy (Bias): How close is the expected value of the MLE to the true parameter values? If the expected value of the MLE is close to the true parameter values, the accuracy is high.
2. Precision: How much does the MLE vary from sample to sample? We measure this by assessing the 95% confidence intervals (BCa, Bias-Corrected and accelerated). If the confidence intervals are both small and have good coverage probability (the proportion of confidence intervals that contain the true parameter values), then the MLE is precise.

Table 2: Weibull Components in Series Configuration

	Shape (k_j)	Scale (λ_j)	MTTF $_j$	$\Pr\{K_i = j\}$	$R_j(\tau; k_j, \lambda_j)$
Component 1	1.2576	994.3661	924.869	0.169	0.744
Component 2	1.1635	908.9458	862.157	0.207	0.698
Component 3	1.1308	840.1141	803.564	0.234	0.667
Component 4	1.1802	940.1342	888.237	0.196	0.711
Component 5	1.2034	923.1631	867.748	0.195	0.711
Series System	NA	NA	222.884	NA	0.175

We begin by specifying the parameters of the series system that will be the central object of our simulation study. We consider the data in Guo et al. (2013), in which they study the reliability of a series system with three components. They fit Weibull components in a series configuration to the data, resulting in an MLE with shape and scale estimates given by the first three components in Table 2. To make the model slightly more complex, we add two more components to this series system, with shape and scale parameters given by the last two components in Table 2. We will refer to this system as the **base** system.

In Section 2.2, we defined a well-designed series system as one that consists of components with similar reliabilities, where we define reliability in two ways, the mean time to failure (MTTF) and the probability that a specific component will be the cause of failure. All things else being equal, components with long MTTFs and with near uniform probability of being the component cause of failure is preferable, otherwise we have a weak link in the system.

The base system defined in Table 2 satisfies this definition of being a well-designed system. We see that there are no components that are significantly less reliable than any of the others, component 1 being the most reliable and component 3 being the least reliable. This is a result of the scales and shapes being similar for each component. In addition, the shapes are larger than 1, which means components are unlikely to fail early.

7.1 Data Generating Process

In this section, we describe the data generating process for our simulation studies. It consists of three parts: the series system, the candidate set model, and the right-censoring model.

Series System Lifetime

We generate data from a Weibull series system with m components. As described in Section 6, the j^{th} component of the i^{th} system has a lifetime distribution given by

$$T_{ij} \sim \text{WEI}(k_j, \lambda_j)$$

and the lifetime of the series system composed of m Weibull components is defined as

$$T_i = \min\{T_{i1}, \dots, T_{im}\}.$$

To generate a data set, we first generate the m component failure times, by efficiently sampling from their respective distributions, and we then set the failure time t_i of the system to the minimum of the component failure times.

Right-Censoring Model

We employ a simple right-censoring model, where the right-censoring time τ is fixed at some known value, e.g., an experiment is run for a fixed amount of time τ , and all systems that have not failed by the end of the experiment are right-censored. The censoring time S_i of the i^{th} system is thus given by

$$S_i = \min\{T_i, \tau\}.$$

So, after we generate the system failure time T_i , we generate the censoring time S_i by taking the minimum of T_i and τ . In our simulation study, we parameterize the right-censoring time τ by the quantile $q = 0.825$ of the series system,

$$\tau = F_{T_i}^{-1}(q).$$

This means that 82.5% of the series systems are expected to fail before time τ and 17.5% of the series are expected to be right-censored. To solve for the 82.5% quantile of the series system, we define the function g as

$$g(\tau) = F_{T_i}(\tau; \boldsymbol{\theta}) - q$$

and find its root using the Newton-Raphson method. See Appendix 10 for the R code that implements this procedure.

Masking Model for Component Cause of Failure

We must generate data that satisfies the masking conditions described in Section 3.1. There are many ways to satisfying the masking conditions. We choose the simplest method, which we call the *Bernoulli candidate set model*. In this model, each non-failed component is included in the candidate set with a fixed probability p , independently of all other components and independently of $\boldsymbol{\theta}$, and the failed component is always included in the candidate set. See Appendix ?? for the R code that implements this model.

7.2 Simulation Scenarios

We define a simulation scenario to be some combination of n (sample size), p (masking probability in our Bernoulli candidate set model), k_3 (shape parameter of the third component), λ_3 (scale parameter of the third component), and q (right-censoring quantile). We are interested in choosing a small number of scenarios that are representative of real-world scenarios and that are interesting to analyze.

Here is an outline of the simulation study for a particular scenario:

1. Fix a combination of simulation parameters to some value, and vary the remaining parameters. For example, if we want to assess how the sampling distribution of the MLE changes with respect to sample size, we might choose some particular values for p , k_3 , λ_3 , and q , and vary the sample size n over the desired range.
2. Simulate $R \geq 300$ datasets from the Data Generating Process (DGP) described in Section 7.1 and compute an MLE for each dataset. We choose R to be large enough so that the sampling distribution of the MLE is well approximated by the empirical distribution of the R MLEs.
3. For each of these R MLEs, compute some function of the MLE, like the BCa confidence intervals or the likelihood ratio test statistic. This will give us R statistics as a Monte-carlo estimate of the sampling distribution of the statistic.
4. Use the R statistics to estimate some property of the sampling distribution of the statistic, e.g., the mean of the MLE or the coverage probability of the BCa confidence intervals, with respect to the parameter(s) we are varying in the scenario, e.g., assess how the coverage probability of the BCa confidence intervals changes with respect to sample size.
5. Visualizze the results and assess the behavior of estimator under the chosen scenario.

For how we run a simulation scenario, see Appendix ??.

7.3 Scenario: Assessing the Impact of Right-Censoring

In this scenario, we use the well-designed series system described in Section 7, Table 2, and we vary the right-censoring quantile (q) from 60% to 100% (no right-censoring), with a component cause of failure masking probability of 21.5% and sample size $n = 100$.




Figure 2: Right-Censoring Quantile vs MLE ($p = 0.215, n = 100$)

Analysis

When a right-censoring event occurs, in order to increase the likelihood of the data, the MLE is nudged in a direction that increases the probability of a right-censoring event at time τ , which is given by $R_{T_i}(t; \theta)$, representing a source of bias in the estimate.

To increase $R_{T_i}(\tau)$, we move in the direction (gradient) of these partial derivatives. The partial derivatives of $R_{T_i}(\tau)$ are given by

$$\begin{aligned}\frac{\partial R_{T_i}(\tau)}{\partial \lambda_j} &= R_{T_i}(\tau; \theta) \left(\frac{\tau}{\lambda_j} \right)^{k_j} \frac{k_j}{\lambda_j}, \\ \frac{\partial R_{T_i}(\tau)}{\partial k_j} &= R_{T_i}(\tau; \theta) \left(\frac{\tau}{\lambda_j} \right)^{k_j} (\log \lambda_j - \log \tau),\end{aligned}$$

for $j = 1, \dots, m$. We see that these partial derivatives are related to the score of a right-censored likelihood contribution in Theorem 6.3. Let us analyze these partial derivatives:

- As τ increases, $R_{T_i}(\tau; \theta)$ decreases, and so the effect right-censoring has on the MLE decreases. This is what we see in Figure 2.
- The partial derivatives with respect to the scale parameters are always positive, so right-censoring positively bias the scale parameter estimates to make right-censoring events more likely. The more right-censoring, the more the positive bias. We see this in Figure 2, where the scale parameter bias decreases as we decrease the probability $(1 - q)$ of a right-censoring event.
- The partial derivative with respect to the shape parameter of the j^{th} component, k_j , is non-negative if $\lambda_j \geq \tau$ and otherwise negative. In our well-designed series system, the scale parameters are large compared to most of the right-censoring times for $\tau(q)$, so the MLE nudges the shape parameter estimates in a positive direction to increase the probability of a right-censoring event $R_{T_i}(\tau)$ at time τ . We see this in Figure 2, where the shape parameter estimates are positively biased for most of the quantiles q .

An alternative way to reason about the effect right-censoring has on the MLE is give by the following. When a right-censoring event occurs, to make the estimates more compatible with the right-censoring event, the estimate is nudged in a direction that increases the MTTF of the components and decreases the probability of component failures before the censoring time τ . In the case of our well-designed system, for most of the right-censoring quatiles q , this is accomplished by nudging the shape and scale parameters in a positive direction, which causes the MLE to be positively biased in the prescence of these censoring times. As the right-censoring quantile q increases, the effect of right-censoring on the MLE decreases.

Here are several key observations:

- *Coverage Probability (CP)*: The CP is well-calibrated, obtaining a value near the nominal 95% level across different right-censoring quantiles. This suggests that the bootstrapped CIs will contain the true value of the parameters with the specified confidence level. The CIs are neither too wide nor too narrow.
- *Dispersion of MLEs*: The shaded regions representing the 95% probability range of the MLEs get narrower as the right-censoring quantile increases. This is an indicator of the increased precision in the estimates as more data is available due to decreased censoring.
- *IQR of Bootstrapped CIs*: The IQR (vertical blue bars) reduces with an increase in sample size. This suggests that the bootstrapped CIs are getting more consistent and focused around a narrower range with larger samples while maintaining a good coverage probability. As we get more data, the bootstrapped CIs are more likely to be closer to each other and the true value of the parameters.

For small right-censoring quantiles (small right-censoring times), they are quite large, but to maintain well-calibrated CIs, this was necessary. The estimator is quite sensitive to the data, and so the bootstrapped CIs are quite wide to account for this sensitivity when the sample contains insufficient information due to censoring.

- *Mean of MLEs*: The red dashed line indicating the mean of MLEs initially is quite biased, but quickly diminishes to negligible levels for scale parameters. The estimates for the shape parameters never reaches zero, but this is potentially due to masking. At a larger sample size, we anticipate the bias in the shape estimates would also decrease to zero.

7.4 Scenario: Assessing the Impact of Sample Size

In this scenario, we use the well-designed series system described in Section 7, Table 2. We fix the masking probability to $p = 0.215$ (moderate masking), we fix the right-censoring quantile to $q = 0.825$ (moderate censoring), and we vary the sample size n from 50 (small sample size) to 1000 (very large sample size).

In Figure 3, we show the effect of the sample size n on the MLEs for the shape and scale parameters. The top four plots only show the effect on the MLEs for the shape and scale parameters of components 1 and 4, since the rest were essentially identical, and the bottom two plots show the coverage probabilities for all parameters.

Here are several key observations:

- *Coverage Probability (CP)*: The CP is well-calibrated, obtaining a value near the nominal 95% level across different sample sizes. This suggests that the bootstrapped CIs will contain the true value of the shape parameter with the specified confidence level. The CIs are neither too wide nor too narrow.
- *Dispersion of MLEs*: The shaded regions representing the 95% probability range of the MLEs get narrower as the sample size increases. This is an indicator of the increased precision in the estimates when provided with more data. This is consistent with the asymptotic properties of the MLE when the regularity conditions are satisfied, e.g., converges in probability to the true value of the parameter as n goes to infinity.
- *IQR of Bootstrapped CIs*: The IQR (vertical blue bars) reduces with an increase in sample size. This suggests that the bootstrapped CIs are getting more consistent and focused around a narrower range with larger samples while maintaining a good coverage probability. As we get more data, the bootstrapped CIs are more likely to be closer to each other and the true value of the scale parameter.

For small sample sizes, they are quite large, but to maintain well-calibrated CIs, this was necessary. The estimator is quite sensitive to the data, and so the bootstrapped CIs are quite wide to account for this sensitivity when the sample size is small and not necessarily representative of the true distribution.

- *Mean of MLEs for Scales*: The red dashed line indicating the mean of MLEs remains stable across different sample sizes and close to the true value, suggesting that the scale MLEs are, on average, reasonably unbiased.
- *Mean of MLEs for Shapes*: The red dashed line is the mean of shape MLEs. Unlike the scale MLEs, we see that for small samples, particularly less than 200, we observe a significant amount of positive bias for shape MLEs. The MLE for the shape parameters in this scenario appear to be more sensitive to the data than the scale parameters.

In Section 7.3, right-censoring was shown to have a very small effect of positively biasing the shape parameter estimates for right-censoring quantiles around $q = 0.825$, but not for the scale parameters. This may partly explain the positive bias in the shape parameter estimates for small sample sizes. However, in Section ??, we show that it may also be related to the masking probability $p = 0.215$. When the sample size is small, the data is not sufficient to overcome the masking effect, and so the MLE is positively biased due to the uncertainty about which component caused the system to fail.

This scenario successfully illustrates the importance of sample size in estimating parameters. The findings align with statistical theory and provide insights into the behavior of these estimators for different sample sizes. In particular, it highlights the sensitivity of the shape parameter estimates to right-censoring and masking for small sample sizes and the importance of having sufficient data to overcome these effects.




Figure 3: Sample Size vs MLEs ($p = 0.215, q = 0.825$)

7.5 Scenario: Assessing the Impact of Masking Probability for Component Cause of Failure

In this scenario, we use the well-designed series system described in Section 7, Table 2. We fix the sample size to $n = 90$ (reasonable sample size) and we fix the right-censoring quantile to $q = 0.825$, and we vary the masking probability from p from 0.1 (very slight masking the component cause of failure) to 0.85 (extreme masking of the component cause of failure).

In Figure 4, we show the effect of the masking probability p on the MLE for the shape and scale parameters. The top four plots only show the effect on the MLEs for the the shape and scale parameters of components 1 and 4, since the rest were essentially identical, and the bottom two plots show the coverage probabilities for all parameters.

Analysis

Here are several key observations:

- *Coverage Probability (CP)*: For the scale parameters, the 95% CI is well-calibrated for masking probabilities up to $p = 0.725$, which is really quite significant, obtaining coverages over 90%, but drops precipitously after that point.

For the shape parameters, the 95% CI is well-calibrated for masking probabilities only up to $p = 0.4$, which is still large, obtaining coverages generally over 90, but begins to drop slowly after that point.

The bootstrapped (BCa) confidence intervals seem fairly well-calibrated for most realistic masking probabilities, constructing CIs that are neither too wide nor too narrow, but when the masking is essentially a fair coin flip, one should probably take the CIs with a grain of salt.

- *Dispersion of MLEs*: The shaded regions representing the 95% quantile of the MLEs become wider as the masking probability increases. This is an indicator of the decreased precision in the estimates when provided with more ambiguous data about the component cause of failure. However, even for fairly significant masking, $p \leq 0.55$, the 95% quantiles are relatively narrow and the CP is well-calibrated, indicating that the MLEs are still relatively precise and accurate.
- *IQR of Bootstrapped CIs*: The IQR (vertical blue bars) show that the bootstrapped BCa CIs are becoming more spread out as the masking probability increases. They are also asymmetric, with the lower bound being more spread out than the upper bound, but this is consistent with the actual behavior of the dispersion of the MLEs, which exhibits the same pattern. The width of the CIs consistently increase as the masking probability increases, which we intuitively expected given the increased uncertainty about the component cause of failure.

After a masking probability of $p = 0.4$, the width of the CIs rapidly increase, which is apparently necessary for the CPs to remain well-calibrated.

- *Mean of MLEs*: The red dashed line indicating the mean of the MLEs remains stable across different masking probabilities, only showing a significant positive bias when the masking probability p becomes quite significant.

7.6 Scenario: Assessing the Impact of Changing the Scale Parameter of a Component

By Equation (??), we see that $MTTF_j$ is proportional to the scale parameter λ_j , which means when we decrease the scale parameter of a component, we proportionally decrease the MTTF. In this scenario, we start with the well-designed series system described in Section 7, Table 2, and we will manipulate the MTTF of component 3 by changing its scale parameter, λ_3 , and observing the effect this has on the MLE. Since the other components had a similar MTTF, we will arbitrarily choose component 1 to represent the other components.




Figure 4: Component Cause of Failure Masking (p) vs MLE

We vary the MTTF of component 3 from around 300 to 1500 and the other components have their MTTFs fixed at around 900, as show in Table 2. We fix the masking probability to $p = 0.215$ (moderate masking), the right-censoring quantile to $q = 0.825$ (moderate censoring), and the sample size to $n = 100$ (moderate sample size).

Analysis

We make two observations about Figure 5:

- When the MTTF of component 3 is much smaller than the other components, the estimate of parameters of component 3 is precise (narrow CIs with high Probability coverage) and accurate (the MLE is close to the true value). This is because component 3 is the component cause of failure in nearly every system failure, and so the data is very informative about the parameters of component 3. Conversely, the estimates of the parameters of the other components is quite poor, with wide CIs and large positive bias. Nonetheless, the coverage probability of the CIs for the other components is still well-calibrated, which means that the CIs will contain the true value of the parameter with a probability around the specified confidence level. So, while we may not have a good point estimates for the parameters, we can still be confident that CIs contain them. That is to say, we have properly quantified our uncertainty about the parameters of the other components.
- When the MTTF of component 3 is much larger than the MTTF of the other components, then component 3 is much less likely to be the component cause of failure, and with a masking probability of $p = 0.215$, it will be in the candidate set with approximately 21.5% probability, but it will generally be a false candidate. The end result is that the estimates of the parameters of component 3 are quite poor, with wide CIs and large positive bias. However, the estimates of the parameters of the other components are quite good, with narrow CIs and small positive bias. The coverage probability of the CIs for the other components are, in comparison, quite good. As the MTTF of component 3 increases and it becomes less likely to be the component cause of failure, the estimates of the parameters of the other components become more precise and accurate.

We also see that the bias is positive for both parameters of component 3. We had not necessarily expected this, but we knew there would be a complex relationship given the presence of right-censoring and masking. When a system is right-censored, or the exact time of failure is observed but the component cause of failure is masked and component 3 is not in the candidate set, then to make component 3 more likely to not be the component cause of failure, its failure rate at that observed time is pushed down and its MTTF is pushed to the right by the MLE. Thus, $\hat{\lambda}_3$ being positively biased is expected. However, k_3 being positively biased is not necessarily expected, but the fact is, decreasing k_3 only has a small impact on the MTTF compared to the scale parameter λ_3 , and the shape parameter may be more particular about when the failures occur. For example, if the shape parameter is large, then the failures may be more likely to occur at the beginning of the lifetime, which would cause the MTTF to be pushed to the right. This is

7.7 Scenario: Assessing the Impact of Changing the Shape Parameter of a Component

We vary the shape parameter of component 3 from 0.1 to 1.9 and observe the effect it has on the MLE and the bootstrapped confidence intervals (BCa). The shape parameter determines the failure characteristics of component 3.

When $k_3 < 1$, this indicates infant mortality, with a decreasing failure rate over time, so even though it has a high failure rate at the beginning of its lifetime, it has a low failure rate at the end of its lifetime and its MTTF is much higher than the other components even though it has a higher probability of failing first. When $k_3 > 1$, this indicates wear-out failures, with an increasing failure rate over time, so even though it has a low failure rate at the beginning of its lifetime, it has a high failure rate at the end of its lifetime and it has a lower probability of failing first.

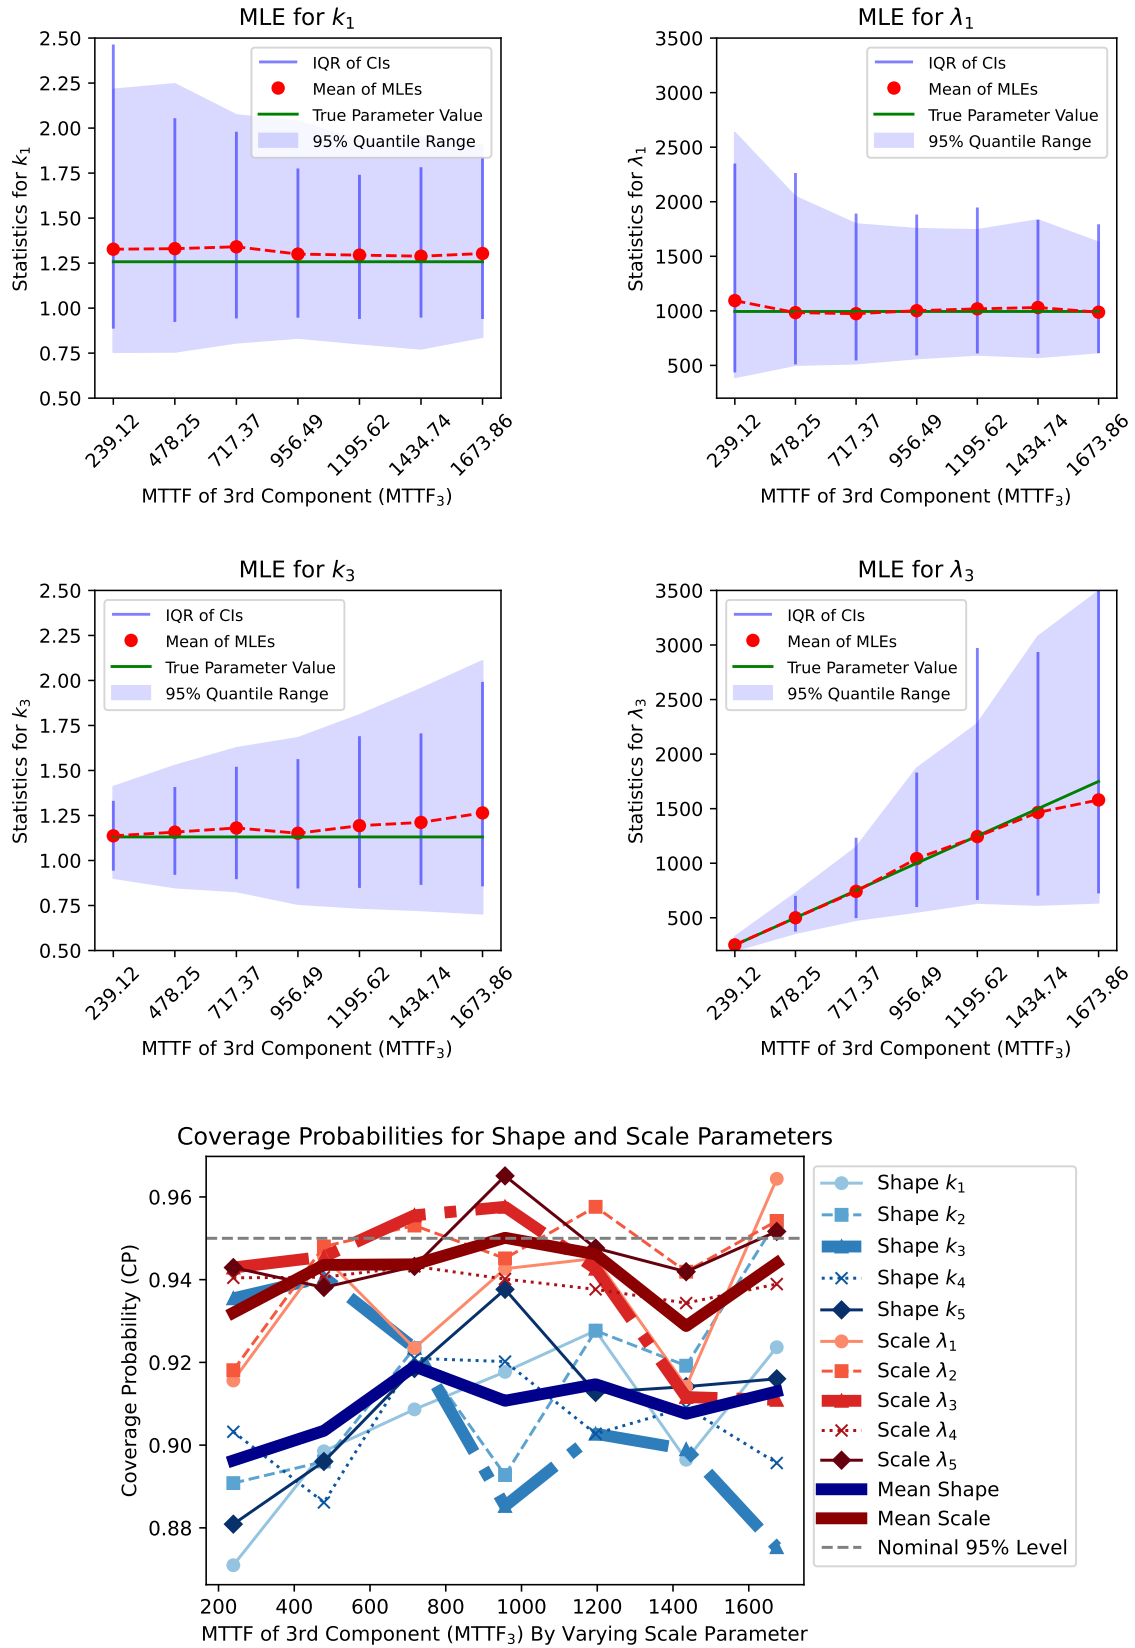


Figure 5: $MTTF_3$ vs MLE By Varying Scale

We analyze the effect of component 3's shape parameter on the MLE and the bootstrapped confidence intervals for the shape and scale parameters of components 1 and 3 (the component we are varying). First, we look at the effect on the scale parameter.

7.7.1 Scales Analysis

In Figure 6, we show the effect of the shape parameter of component 3 on the MLE and the bootstrapped confidence intervals for the shape parameters of components 1 and 3. We see that the mean MLE, in green, is relatively close to the true value, in red, for the scale parameter of both components. There is a slight positive bias, which may be due to the fact that the data is right-censored with moderate masking of the component cause of failure. We see that as the probability of component 3 being the cause of failure increases, the bootstrapped confidence intervals generally increase in width, with the exception of when $k_3 < 1$ which causes $\Pr\{K_i = 3\}$ to be very small and as $\Pr\{K_i = 3\}$ approaches 0.2, all of the components are approximately equally likely to be the component cause of failure, and so the CIs seem to be fairly small for all scale parameters.

However, for $\Pr\{K_i = 3\} > 0.5$, we see that the the mean MLE begins to increase significantly for λ_3 . This is somewhat unexpecte; we might think that, because its probability of being the component cause of failure is higher, that we would estimate λ_3 to be lower to proportionately decrease its MTTF. However, the fact is that the shape parameter has a much bigger impact.

Also, the coverage probabilities of the confidence intervals for the scale parameters decreases for the scale parameter of components other than 3 as $\Pr\{K_i = 3\}$ increases, but the coverage probability for the scale parameter of component 3 increases. This may be because, as $\Pr\{K_i = 3\}$ increases, we are more likely to observe a failure of component 3, and so we have more information about its parameters and are able to estimate them more accurately.

7.7.2 Shapes Analysis

Now, we look at the effect of the shape parameter of component 3 on the MLE and the bootstrapped confidence intervals for the shape parameters of components 1 and 3.

In Figure 6, we show the effect of the shape parameter of component 3 on the MLE and the bootstrapped confidence intervals for the shape parameters of components 1 and 3.

We see that the bias for k_1 slowly increases (positive bias) as $\Pr\{K_i = 3\}$ increases, and the bias for k_3 slowly decreases (positive bias) to 0 as $\Pr\{K_i = 3\}$ increases. This makes sense, as a larger positive bias for k_1 means that the MLE is nudging the shape parameter of component 1 to be larger so that component 1 is less likely to be the cause of failure. Similarly, a smaller positive bias for k_3 means that the MLE is nudging the shape parameter of component 3 to be smaller so that component 3 is more likely to be the cause of failure.

The confidence intervals for k_1 also become quite wide as $\Pr\{K_i = 3\}$ increases, which is expected since we observe fewer failures of component 1 as $\Pr\{K_i = 3\}$ increases, and so we have less information about its parameters and are less able to estimate them accurately. Conversely, the confidence intervals for k_3 become narrower as $\Pr\{K_i = 3\}$ increases, which is also expected, since we observe more failures of component 3 as $\Pr\{K_i = 3\}$ increases, and so we have more information about its parameters and are more able to estimate them accurately. The CI widths for k_3 becomes extremely small for $\Pr\{K_i = 3\} > 0.3$.

The coverage probabilities are generally less well-calibrated for the shape parameters compared to the scale parameters, but they are still reasonably well-calibrated for $\Pr\{K_i = 3\} < 0.4$. For k_3 , the coverage probabilities are very well-calibrated for all values of $\Pr\{K_i = 3\}$, but improve as $\Pr\{K_i = 3\}$ increases due to the fact that we observe more failures of component 3 as $\Pr\{K_i = 3\}$ increases and thus have more informationa bout k_3 for our estimate.

7.8 Implications and Recommendations

Shape parameters The shape parameters are more sensitive to the data than the scale parameters. The coverage probabilities were also consistently less well-calibrated for the shape parameters compared to the scale parameters, but they were still reasonably well-calibrated for most scenarios, with a lower bound of 85%

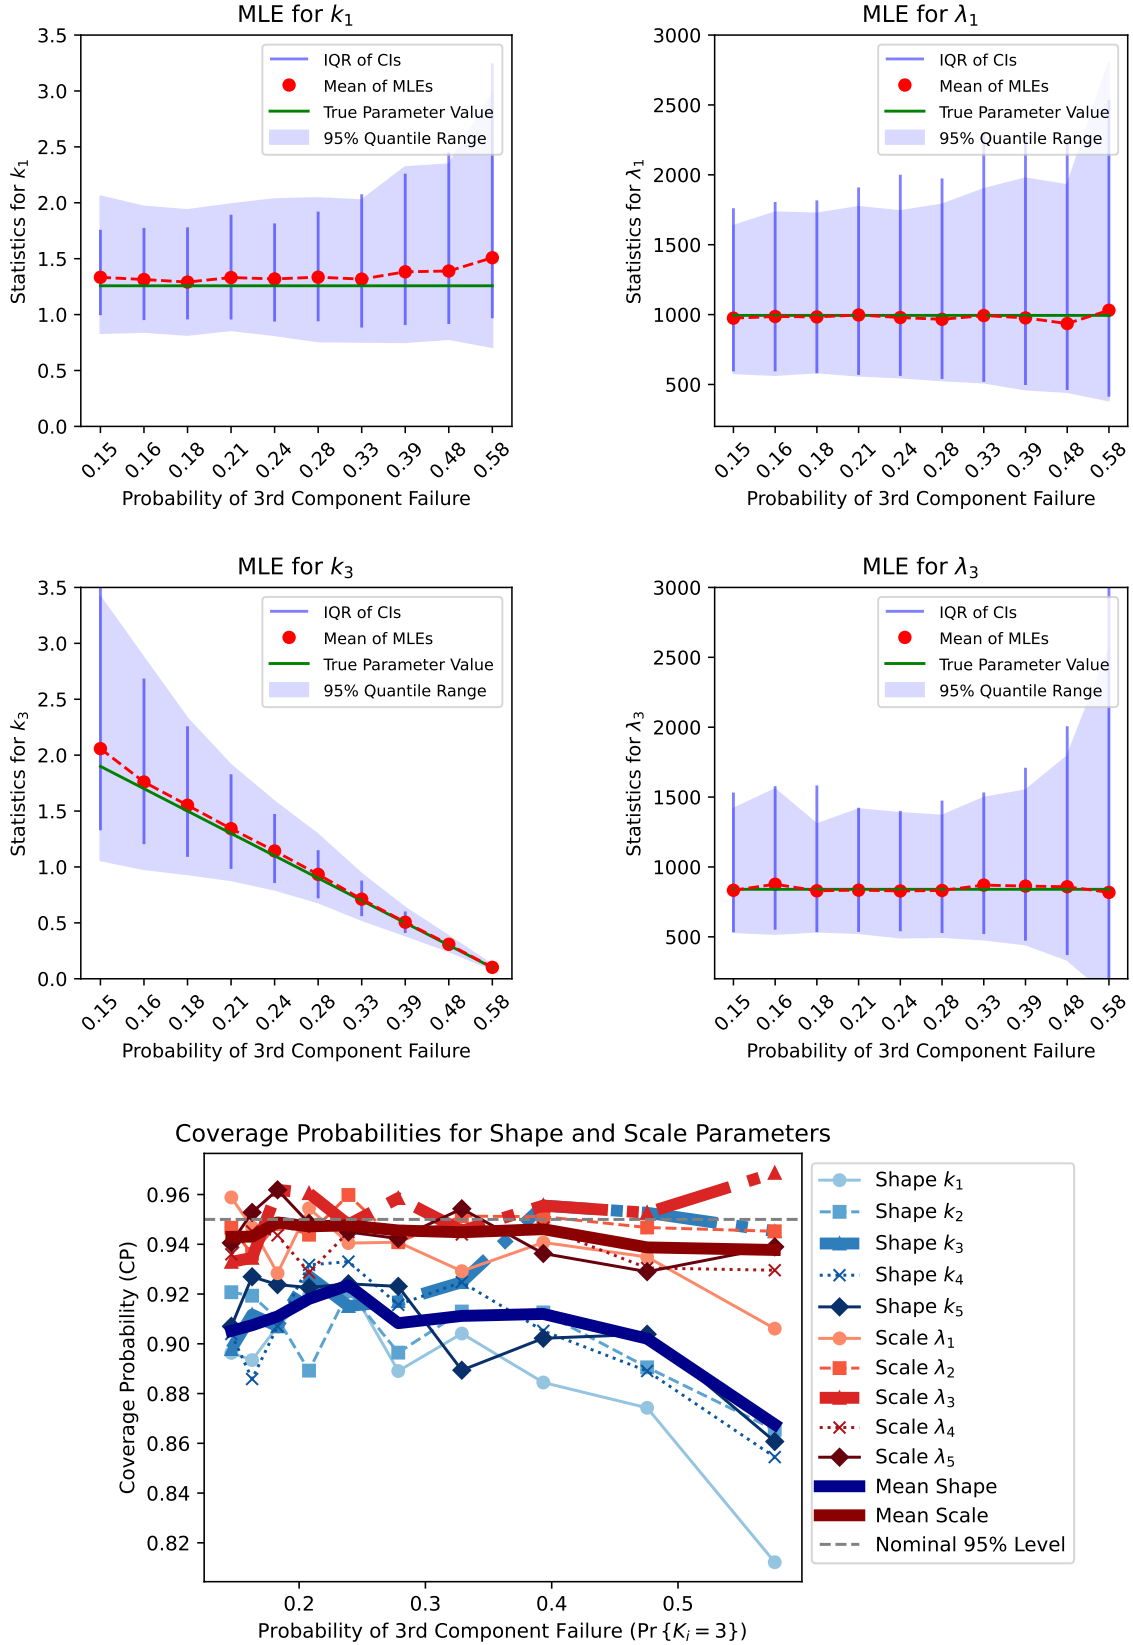


Figure 6: Probability of Component 3 Failure vs MLE

for the most extreme scenarios, e.g., Bernoulli masking probabilities of $p = 0.7$ with a relatively small sample ($n = 100$) and moderate right-censoring ($q = 0.825$).

The shape parameters define the failure characteristics of the components, but they have a complex relationship to the MTTF and other reliability measures. Ideally, the shape parameters will all be greater than 1, indicating wear-out failures, but this is not always the case, and sometimes it may be the case that some components have a shape parameter greater than 1 while others have a shape parameter less than 1. Due to the sensitivity of the shape parameters to the data, and the fact that the shape parameters are less intuitive than the scale parameters, we recommend that the shape parameters be interpreted with caution, particularly when the sample size is small.

Since a well-designed system has components with similar failure characteristics, the shape parameters should be reasonably aligned. To mitigate the problems with the shape parameters, both from an analytic and variability perspective, we may be motivated to use the reduced model that assumes homogeneity in the shape parameters, as described in Section 6.3. The reduced model offers interpretability (the series system is itself Weibull) and reduced estimator variability (only $m + 1$ parameters instead of $2m$), but it must adequately describe the data.

Scale parameters The scale parameters are less sensitive to the data than the shape parameters. The coverage probabilities were also consistently well-calibrated for all scenarios, with a lower-bound of around 90% for the most extreme scenarios. The scale parameters are linearly related to the MTTF, and have a simpler relationship to the other reliability measures too, which makes it easier to interpret and understand.

Sample Size The MLE is very sensitive to the data, and so the bootstrapped CIs are quite wide to account for this sensitivity when the sample size is small and not necessarily representative of the true distribution.

Right-Censoring For our well-designed system, the mean MLE for the shape parameters were mildly sensitive to right-censoring, but even significant censoring around 50% only introduced a small positive bias. The mean MLE for the scale parameters were not very sensitive to right-censoring.

The sampling variance of the MLE for both the shape and scale parameters were somewhat sensitive to right-censoring, as evidenced by the decreasing width of the confidence intervals as a function of increasing right-censoring times. However, the coverage probabilities of the confidence intervals were well-calibrated for the scale parameters

the bootstrapped CIs are quite wide to account for this sensitivity when the sample contains insufficient information due to censoring. The

- The BCa (bootstrapped) confidence intervals are well-calibrated in all but the most extreme circumstances, constructing CIs that are neither too wide nor too narrow, with two notable exceptions:
 - When the masking probability is essentially a fair coin flip, one should probably take the CIs with a grain of salt.
 - When the components have radically differentiated failure characteristics, the CIs for the parameters of the other components are quite wide, but the coverage probability is still well-calibrated.

8 Simulation Study: Full Model vs Reduced Model

In Section 6, we developed the full model with Weibull components and we conducted a sensitivity analysis of its MLE in Section 7. Here, our focus shifts to a sensitivity analysis aimed at understanding when it is appropriate to use the reduced model that assumes homogeneity in the shape parameters of the components, as described in Section 6.3. The reduced model offers interpretability (the series system is itself Weibull) and reduced estimator variability (only $m + 1$ parameters instead of $2m$), but it must adequately describe the data.

We aim to assess the appropriateness of the reduced model under varying sample sizes and shape parameters of the third component (k_3). We employ a simulation study using the likelihood ratio test (LRT) for this purpose, where the null hypothesis, H_0 , assumes homogenous shape parameters.

We take the well-designed series system described in Section 7, Table 2, and manipulate the shape parameter of the third component (k_3) to cause the components to have different failure characteristics. Recall that $k_3 = 1.1308$ corresponds to a *well-designed* series system, where component shapes are reasonably aligned. We also vary the sample size n to assess the impact of sample size on the appropriateness of the reduced model.

Figure 7 provides a contour plot with varying sample sizes along the x -axis, shape of component 3 along the y -axis, and median p -value along the color scale. The contour lines corresponding to p -values of 0.05 and 0.1 are often used as a threshold for statistical significance. Points outside of these contours in dark blue are indicative of significant evidence against the null model and points inside of these contours in light blue for 0.1 and green for 0.05 are indicative of insufficient evidence against the null model.

Figure 8 provides a plot of the median p -value against the sample size for the well-designed system, where the shape parameter of component 3 is fixed at 1.1308. The 95th percentile of the p -values is also provided as a more stringent criterion for statistical significance.




Figure 7: p -Value vs Sample Size and Shape k_3




Figure 8: p -Value vs Sample Size for Well-Designed System

Sensitivity to Sample Size (n)

- The sample size is an essential aspect of hypothesis testing, as it affects the test's power, which is the probability of correctly rejecting the null hypothesis when it is false. In the contour plot in Figure 7, as n increases, the contours trend lower. This indicates that larger samples result in smaller median p -values, implying that the power of the LRT increases with the sample size. However, its power is quite low for small samples, particularly for values of k_3 somewhat close to the shape parameters of the other components in the system.
- Recall that in the well-designed series system, $k_3 = 1.1308$. In this case, even very large sample sizes do not produce evidence against the null model, indicating robust compatibility.
- In Figure 8, we fix k_3 at 1.1308 and vary the sample size. The median p -value only manages to drop below the 0.05 threshold with sample sizes around 10000. In the more stringent criterion given by the 95th percentile of the p -values, nearly 30000 observations are necessary to reject the null hypothesis in 95% of the simulations.

Sensitivity to Shape Parameter (k_3)

- In Figure 7, for a given shape parameter, increasing the sample size tends to decrease the median p -value. Larger samples provide more information about the parameters, which increases the power of the LRT.
- The median p -values in the vicinity of the line $k_3 = 1.1308$ are high across various sample sizes, indicating that the null model is a good fit. As k_3 deviates from this line, the median p -value diminishes, indicating increasing evidence against the null model.

8.1 Implications and Recommendations

The power of the test for a well-designed series system is quite low, requiring many thousands of observations before the test has sufficient power to reject the null hypothesis. But, this is not necessarily a bad thing. The reduced model is quite simple and interpretable, and is by definition a good fit for a well-designed series system.

The findings suggest that the reduced model is particularly apt when the system is well-designed, even for very large samples. Practitioners should weigh the trade-offs between the simplicity of the reduced model and the adequacy in describing the data, with consideration of the available sample size and the characteristics of the system being modeled.

For systems believed to be well-designed, employing the null model is supported both statistically and practically due to its simplicity, reduced estimator variability, and analytical tractability.

In the absence of prior information, or if the shape parameter significantly diverges from the well-designed value, the choice between models should be undertaken with caution. More complex models may be favorable, especially with large sample sizes.

9 Future Work

Relaxation of Conditions Investigate relaxations of Conditions 1, 2, and 3. Condition 1 stipulates that the failed component is always in the candidate set,

$$\Pr\{K_i \in \mathcal{C}_i\} = 1.$$

Instead, we could model this as a probability, where the probability of the failed component being in the candidate set is a function of the failure time T_i and the component cause of failure K_i ,

$$\Pr\{K_i \in \mathcal{C}_i | K_i = j, T_i = t_i\} = g(j, t_i).$$

Condition 2 stipulates that

$$\Pr\{\mathcal{C}_i = c_i | T_i = t_i, K_i = j\} = \Pr\{\mathcal{C}_i = c_i | T_i = t_i, K_i = j'\}$$

for all $j, j' \in c_i$. We call this an *uninformed* candidate set, since the conditional probability of the candidate set given the failure time and component cause of failure is independent of the component cause of failure. We could relax this condition to allow for *informed* candidate sets, where the conditional probability of the candidate set given the failure time and component cause of failure is dependent on the component cause of failure.

In each of these violations or relaxations, we can either construct a new likelihood model that takes this relaxation into account, or we can use the existing likelihood model and assess the sensitivity of the estimator to this violation. A potentially interesting way to do the latter is by using KL-divergence to measure the distance between, for instance, the uninformed and informed candidate set models, and then assess the sensitivity of the estimators to this distance.

Expanded Sensitivity Analyses

- Assess the sensitivity of the reduced model to the presence of masking and right-censoring. Since it has fewer parameters, it may be more robust to these effects.
- Assess the sensitivity of the full and reduced models to the number (m) of components in the series system.

Data Augmentation Assess the robustness of Data Augmentation (DA) as an implicit prior, particularly for the reduced model, where we may adopt the prior that the system is well-designed. By incorporating DA into the procedure, we may be able to significantly improve the performance of the estimator on extremely small samples.

Unlike a full Bayesian approach, where we would need to specify a prior for the parameters, DA is an implicit prior that need not be explicitly specified. It is a form of regularization that improves the performance of the estimator by leveraging the structure of the model and the data. It is particularly useful for small samples, where the data is not sufficient to overcome the masking and censoring effects.

Penalized Likelihood Assess the use of penalized likelihood methods instead of DA as a form of regularization. For instance, we can add a penalty term to the log-likelihood function that penalizes the likelihood when the shape parameters are not close to each other. Instead of simply using a reduced model, we can use a penalized likelihood approach to encourage the shape parameters to be close to each other, but not necessarily equal.

10 Conclusion

Section 8 explored the circumstances under which the reduced model is appropriate through a sensitivity analysis involving sample size and the shape parameter of component 3. The insights gained from this section were particularly relevant for practical applications where model interpretability is important and the samples are limited.

Appendix A: R Code For Log-likelihood Function

The following code is the log-likelihood function for the Weibull series system with a likelihood model that includes masked component cause of failure and right-censoring. It is implemented in the R library `wei.series.md.c1.c2.c3` and is available on GitHub.

For clarity and brevity, we removed some of the functionality that is not relevant to the analysis in this paper.

```
## Generates a log-likelihood function for a Weibull series system with respect
## to parameter `theta` (shape, scale) for masked data with candidate sets
## that satisfy conditions C1, C2, and C3 and right-censored data.
##
## @param df (masked) data frame
## @param theta parameter vector (shape1, scale1, ..., shapem, scalem)
## @returns Log-likelihood with respect to `theta` given `df`
loglik_wei_series_md_c1_c2_c3 <- function(df, theta) {
  n <- nrow(df)
  C <- md_decode_matrix(df, candset)
  m <- ncol(C)
  delta <- df[[right_censoring_indicator]]
  t <- df[[lifetime]]
  k <- length(theta)
  shapes <- theta[seq(1, k, 2)]
```

```

scales <- theta[seq(2, k, 2)]

s <- 0
for (i in 1:n) {
  s <- s - sum((t[i] / scales)^shapes)
  if (delta[i]) {
    s <- s + log(sum(shapes[C[i, ]] / scales[C[i, ]] *
      (t[i] / scales[C[i, ]])^(shapes[C[i, ]] - 1)))
  }
}
s
}

```

Appendix B: R Code For Score Function

The following code is the score function (gradient of the log-likelihood function with respect to θ) for the Weibull series system with a likelihood model that includes masked component cause of failure and right-censoring. It is implemented in the R library `wei.series.md.c1.c2.c3` and is available on GitHub.

For clarity and brevity, we removed some of the functionality that is not relevant to the analysis in this paper.

```

## Computes the score function (gradient of the log-likelihood function) for a
## Weibull series system with respect to parameter `theta` (shape, scale) for masked
## data with candidate sets that satisfy conditions C1, C2, and C3 and right-censored
## data.
##
## @param df (masked) data frame
## @param theta parameter vector (shape1, scale1, ..., shapem, scalem)
## @returns Score with respect to `theta` given `df`
score_wei_series_md_c1_c2_c3 <- function(df, theta) {
  n <- nrow(df)
  C <- md_decode_matrix(df, candset)
  m <- ncol(C)
  delta <- df[[right_censoring_indicator]]
  t <- df[[lifetime]]
  shapes <- theta[seq(1, length(theta), 2)]
  scales <- theta[seq(2, length(theta), 2)]
  shape_scores <- rep(0, m)
  scale_scores <- rep(0, m)

  for (i in 1:n) {
    rt.term.shapes <- -(t[i] / scales)^shapes * log(t[i] / scales)
    rt.term.scales <- (shapes / scales) * (t[i] / scales)^shapes

    # Initialize mask terms to 0
    mask.term.shapes <- rep(0, m)
    mask.term.scales <- rep(0, m)

    if (delta[i]) {
      cindex <- C[i, ]
      denom <- sum(shapes[cindex] / scales[cindex] * (t[i] / scales[cindex])^(shapes[cindex] - 1))

      numer.shapes <- 1 / t[i] * (t[i] / scales[cindex])^shapes[cindex] *

```

```

        (1 + shapes[cindex] * log(t[i] / scales[cindex]))
mask.term.shapes[cindex] <- numer.shapes / denom

    numer.scales <- (shapes[cindex] / scales[cindex])^2 * (t[i] / scales[cindex])^(shapes[cindex])
    mask.term.scales[cindex] <- numer.scales / denom
}

shape_scores <- shape_scores + rt.term.shapes + mask.term.shapes
scale_scores <- scale_scores + rt.term.scales - mask.term.scales
}

scr <- rep(0, length(theta))
scr[seq(1, length(theta), 2)] <- shape_scores
scr[seq(2, length(theta), 2)] <- scale_scores
scr
}

```

Appendix C: R Code For Simulation of Scenarios

The following code is the Monte-carlo simulation code for running the scenarios described in Section 7.

```

#### Setup simulation parameters here ####
theta <- c(
  shape1 = 1.2576, scale1 = 994.3661,
  shape2 = 1.1635, scale2 = 908.9458,
  shape3 = NA, scale3 = 840.1141,
  shape4 = 1.1802, scale4 = 940.1342,
  shape5 = 1.2034, scale5 = 923.1631
)

shapes3 <- c(1.1308) # shape 3 true parameter values to simulate
scales3 <- c(840.1141) # scale 3 true parameter values to simulate
N <- c(30, 60, 100) # sample sizes to simulate
P <- c(.215) # masking probabilities to simulate
Q <- c(.825) # right censoring probabilities to simulate
R <- 1000L # number of simulations per scenario
B <- 1000L # number of bootstrap samples
max_iter <- 125L # max iterations for MLE
max_boot_iter <- 125L # max iterations for bootstrap MLE
n_cores <- detectCores() - 1 # number of cores to use for parallel processing
filename <- "data" # filename prefix for output files

#### Simulation code below here ####
library(tidyverse)
library(parallel)
library(boot)
library(algebraic.mle) # for `mle_boot`
library(wei.series.md.c1.c2.c3) # for `mle_lbfgsb_wei_series_md_c1_c2_c3` etc

file.meta <- paste0(filename, ".txt")
file.csv <- paste0(filename, ".csv")
if (file.exists(file.meta)) {
  stop("File already exists: ", file.meta)
}

```

```

}
if (file.exists(file.csv)) {
  stop("File already exists: ", file.csv)
}

shapes <- theta[seq(1, length(theta), 2)]
scales <- theta[seq(2, length(theta), 2)]
m <- length(shapes)

sink(file.meta)
cat("bootstrap of confidence intervals:\n")
cat("  simulated on: ", Sys.time(), "\n")
cat("  type: ", ci_method, "\n")
cat("weibull series system:\n")
cat("  number of components: ", m, "\n")
cat("  scale parameters: ", scales, "\n")
cat("  shape parameters: ", shapes, "\n")
cat("simulation parameters:\n")
cat("  shapes3: ", shapes3, "\n")
cat("  scales3: ", scales3, "\n")
cat("  N: ", N, "\n")
cat("  P: ", P, "\n")
cat("  Q: ", Q, "\n")
cat("  R: ", R, "\n")
cat("  B: ", B, "\n")
cat("  max_iter: ", max_iter, "\n")
cat("  max_boot_iter: ", max_boot_iter, "\n")
cat("  n_cores: ", n_cores, "\n")
sink()

for (scale3 in scales3) {
  for (shape3 in shapes3) {
    for (n in N) {
      for (p in P) {
        for (q in Q) {
          shapes[3] <- shape3
          theta["shape3"] <- shape3

          cat("[starting scenario: scale3 = ", scale3, ",
              shape3 = ", shape3, ", n = ", n, ", p = ", p, ", q = ", q, "]\n")
          tau <- qwei_series(p = q, scales = scales, shapes = shapes)

          # we compute R MLEs for each scenario
          shapes.mle <- matrix(NA, nrow = R, ncol = m)
          scales.mle <- matrix(NA, nrow = R, ncol = m)
          shapes.lower <- matrix(NA, nrow = R, ncol = m)
          shapes.upper <- matrix(NA, nrow = R, ncol = m)
          scales.lower <- matrix(NA, nrow = R, ncol = m)
          scales.upper <- matrix(NA, nrow = R, ncol = m)

          iter <- 0L
          repeat {
            retry <- FALSE

```

```

tryCatch(
  {
    repeat {
      df <- generate_guo_weibull_table_2_data(
        shapes = shapes, scales = scales, n = n, p = p, tau = tau
      )

      sol <- mle_lbfgsb_wei_series_md_c1_c2_c3(
        theta0 = theta, df = df, hessian = FALSE,
        control = list(maxit = max_iter, parscale = theta)
      )
      if (sol$convergence == 0) {
        break
      }
      cat("[", iter, "] MLE did not converge, retrying.\n")
    }

    mle_solver <- function(df, i) {
      mle_lbfgsb_wei_series_md_c1_c2_c3(
        theta0 = sol$par, df = df[i, ], hessian = FALSE,
        control = list(maxit = max_boot_iter, parscale = sol$par)
      )$par
    }

    # do the non-parametric bootstrap
    sol.boot <- boot(df, mle_solver, R = B, parallel = "multicore", ncpus =
  ),
  error = function(e) {
    cat("[error] ", conditionMessage(e), "\n")
    cat("[retrying scenario: n = ", n, ", p = ", p, ", q = ", q, "\n")
    retry <-> TRUE
  }
)
if (retry) {
  next
}
iter <- iter + 1L
shapes.mle[iter, ] <- sol$par[seq(1, length(theta), 2)]
scales.mle[iter, ] <- sol$par[seq(2, length(theta), 2)]

tryCatch(
  {
    ci <- confint(mle_boot(sol.boot), type = ci_method, level = ci_level)
    shapes.ci <- ci[seq(1, length(theta), 2), ]
    scales.ci <- ci[seq(2, length(theta), 2), ]
    shapes.lower[iter, ] <- shapes.ci[, 1]
    shapes.upper[iter, ] <- shapes.ci[, 2]
    scales.lower[iter, ] <- scales.ci[, 1]
    scales.upper[iter, ] <- scales.ci[, 2]
  },
  error = function(e) {
    cat("[error] ", conditionMessage(e), "\n")
  }
)

```

```
)  
if (iter %% 5 == 0) {  
  cat("[iteration ", iter, "] shapes = ", shapes.mle[iter, ], "scales = ", scales.mle[iter, ], "\n")  
}  
  
if (iter == R) {  
  break  
}  
}  
  
df <- data.frame(  
  n = rep(n, R), rep(scale3, R), rep(shape3, R),  
  p = rep(p, R), q = rep(q, R), tau = rep(tau, R), B = rep(B, R),  
  shapes = shapes.mle, scales = scales.mle,  
  shapes.lower = shapes.lower, shapes.upper = shapes.upper,  
  scales.lower = scales.lower, scales.upper = scales.upper  
)  
  
write.table(df,  
  file = file.csv, sep = ",", row.names = FALSE,  
  col.names = !file.exists(file.csv), append = TRUE  
)  
}  
}  
}
```

Appendix D: Bernoulli Candidate Set Model

```

#' Bernoulli candidate set model is a particular type of *uninformed* model.
#' This model satisfies conditions C1, C2, and C3.
#' The failed component will be in the corresponding candidate set with
#' probability 1, and the remaining components will be in the candidate set
#' with probability `p` (the same probability for each component). `p`
#' may be different for each system, but it is assumed to be the same for
#' each component within a system, so `p` can be a vector such that the
#' length of `p` is the number of systems in the data set (with recycling
#' if necessary).
#'
#' @param df masked data.
#' @param p a vector of probabilities ( $p[j]$  is the probability that the  $j$ th
#' system will include a non-failed component in its candidate set,
#' assuming the  $j$ th system is not right-censored).
md_bernoulli_cand_c1_c2_c3 <- function(df, p) {
  n <- nrow(df)
  p <- rep(p, length.out = n)
  Tm <- md_decode_matrix(df, comp)
  m <- ncol(Tm)
  Q <- matrix(p, nrow = n, ncol = m)
  Q[cbind(1:n, apply(Tm, 1, which.min))] <- 1
  Q[!df[[right_censoring_indicator]], ] <- 0
}

```

```

df %>% bind_cols(md_encode_matrix(Q, prob))
}

```

Appendix E: Series System Quantile Function

```

#' Quantile function (inverse of the cdf).
#' By definition, the quantile `p` * 100% for a strictly monotonically increasing
#' cdf `F` is the value `t` that satisfies `F(t) - p = 0`.
#' We solve for `t` using newton's method.
#'
#' @param p vector of probabilities.
#' @param shapes vector of weibull shape parameters for weibull lifetime
#'             components
#' @param scales vector of weibull scale parameters for weibull lifetime
#'             components
qwei_series <- function(p, shapes, scales) {
  t0 <- 1
  repeat {
    t1 <- t0 - sum((t0/scales)^shapes) + log(1-p) /
      sum(shapes*t0^(shapes-1)/scales^shapes)
    if (abs(t1 - t0) < tol) {
      break
    }
    t0 <- t1
  }
  return(t1)
}

```

Appendix F: Likelihood Ratio Test

In order to determine if a reduced model (e.g., Weibull series system in which all of the shape parameters are homogeneous) is appropriate, a hypothesis test may be conducted to determine if there is statistically significant evidence in support of the null hypothesis H_0 , e.g., that all of the shape parameters are homogeneous.

Given that we employ a well-defined likelihood model, the likelihood ratio test (LRT) is a good choice. The LRT statistic is given by

$$\Lambda = -2(\ell_R - \ell_F)$$

where ℓ_R is the log-likelihood of the null (reduced) model (the log-likelihood of the reduced model evaluated at its MLE given a random sample) and ℓ_F is the log-likelihood of the full model. Under the null model, the LRT statistic is asymptotically distributed chi-squared with $m - 1$ degrees of freedom, where m is the number of components in the series system,

$$\Lambda \sim \chi_{m-1}^2.$$

If the LRT statistic is greater than the critical value of the chi-squared distribution with $m - 1$ degrees of freedom, $\chi_{m-1, 1-\alpha}^2$, where α denotes the significance level, then we find the data to be incompatible with the null hypothesis H_0 .

References

Abernethy, R. B. (2006). *New Weibull Handbook*. Abernethy, 5th edition.

- Bain, L. and Engelhardt, M. (1992). *Introduction to Probability and Mathematical Statistics*. Duxbury Press, second edition.
- Casella, G. and Berger, R. L. (2002). *Statistical Inference*. Duxbury Advanced Series.
- Efron, B. (1987). Better bootstrap confidence intervals. *Journal of the American Statistical Association*, 82(397):171–185.
- Efron, B. and Tibshirani, R. J. (1994). *An introduction to the bootstrap*. CRC press.
- Guess, F. M., Hodgson, T. J., and Usher, J. S. (1991). Estimating system and component reliabilities under partial information on cause of failure. *Journal of Statistical Planning and Inference*, 29:75–85.
- Guo, H., Niu, P., and Szidarovszky, F. (2013). Estimating component reliabilities from incomplete system failure data. *Annual Reliability and Maintainability Symposium (RAMS)*, pages 1–6.
- Klein, J. P. and Moeschberger, M. L. (2005). *Survival analysis: techniques for censored and truncated data*. Springer Science & Business Media.
- Lehmann, E. L. and Casella, G. (1998). *Theory of Point Estimation*. Springer Science & Business Media.
- Taleb, N. N. (2007). *The Black Swan: The Impact of the Highly Improbable*. Random House.
- Wu, C. F. J. (1983). On the convergence properties of the em algorithm. *The Annals of Statistics*, 11(1):95–103.